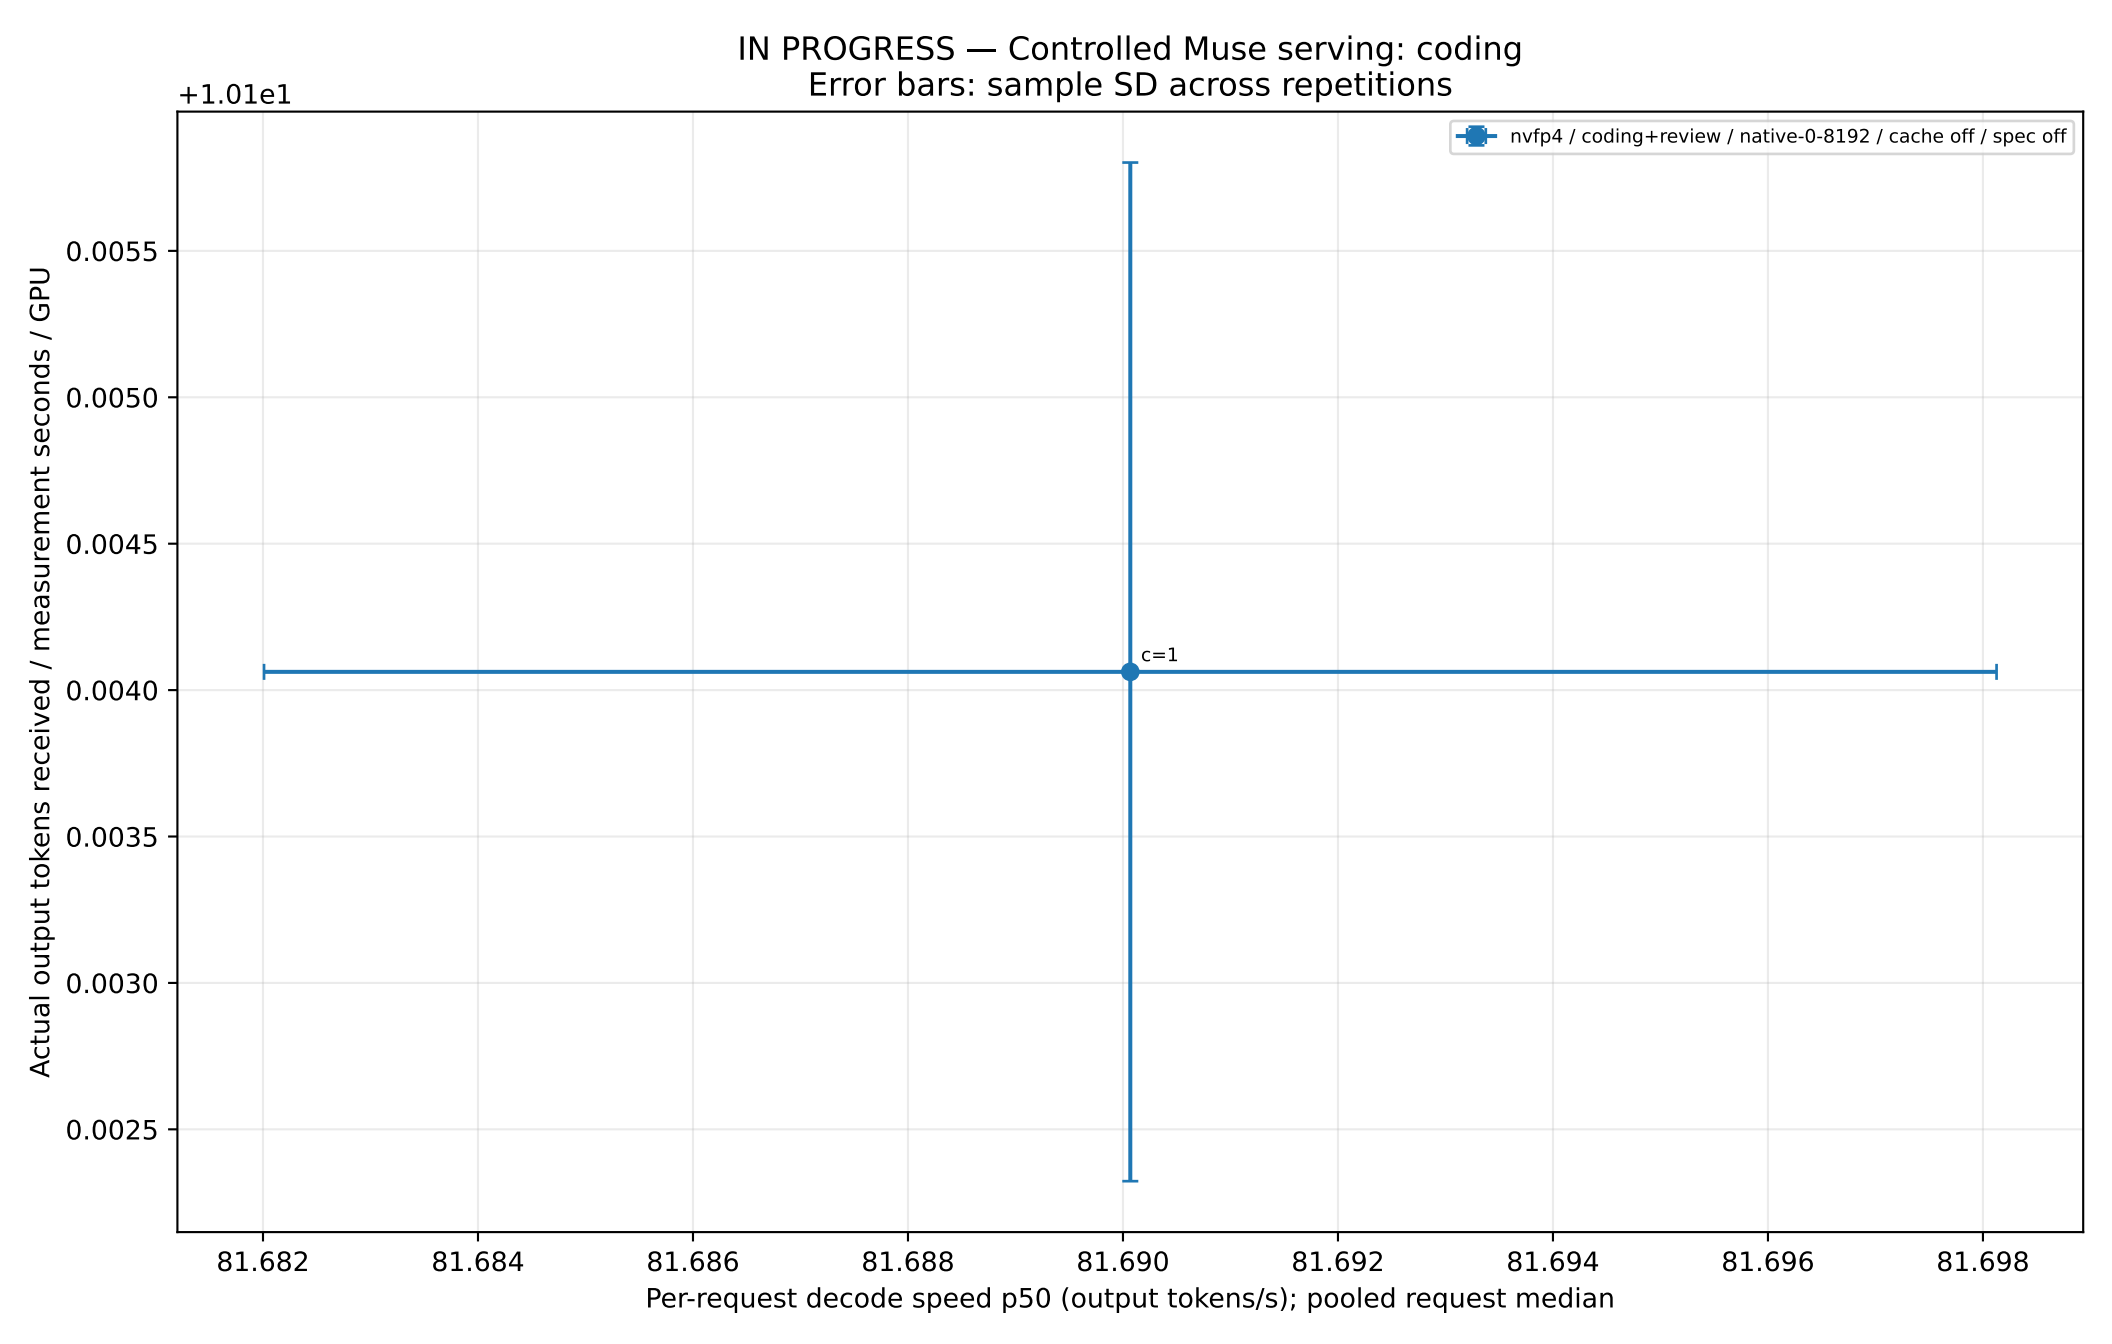
IN PROGRESS — Controlled Muse serving: coding
Error bars: sample SD across repetitions
+1.01e1
nvfp4 / coding+review / native-0-8192 / cache off / spec off
c=1
Actual output tokens received / measurement seconds / GPU
0.0055
0.0050
0.0045
0.0040
0.0035
0.0030
0.0025
81.682
81.684
81.686
81.688
81.690
81.692
81.694
81.696
81.698
Per-request decode speed p50 (output tokens/s); pooled request median