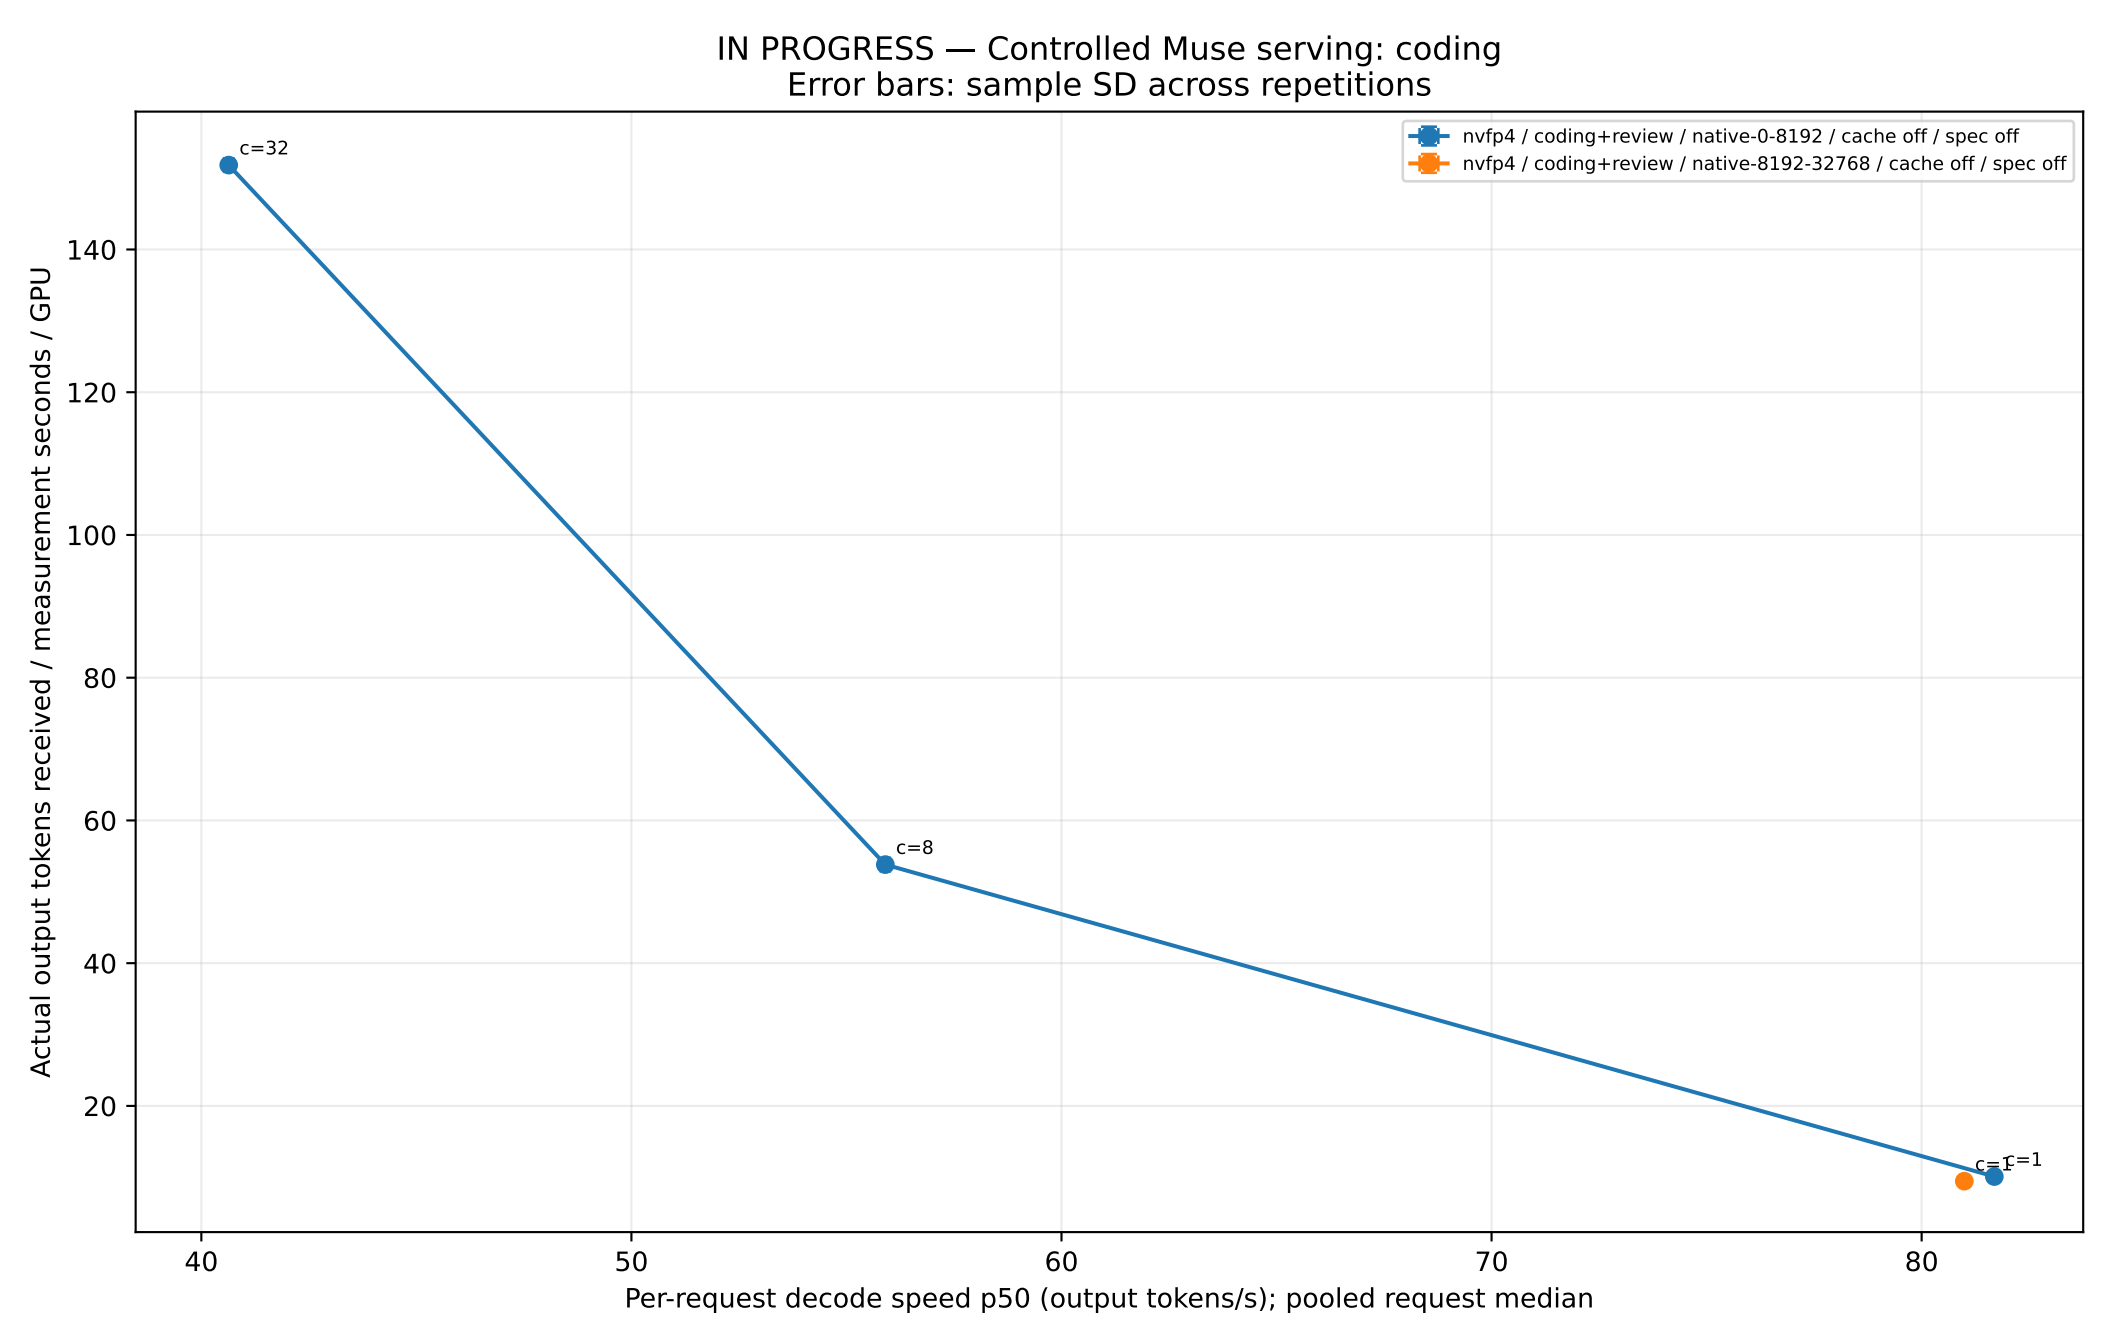
IN PROGRESS — Controlled Muse serving: coding
Error bars: sample SD across repetitions
nvfp4 / coding+review / native-0-8192 / cache off / spec off
nvfp4 / coding+review / native-8192-32768 / cache off / spec off
c=32
c=8
c=1
c=1
Actual output tokens received / measurement seconds / GPU
Per-request decode speed p50 (output tokens/s); pooled request median
20
40
60
80
100
120
140
40
50
60
70
80