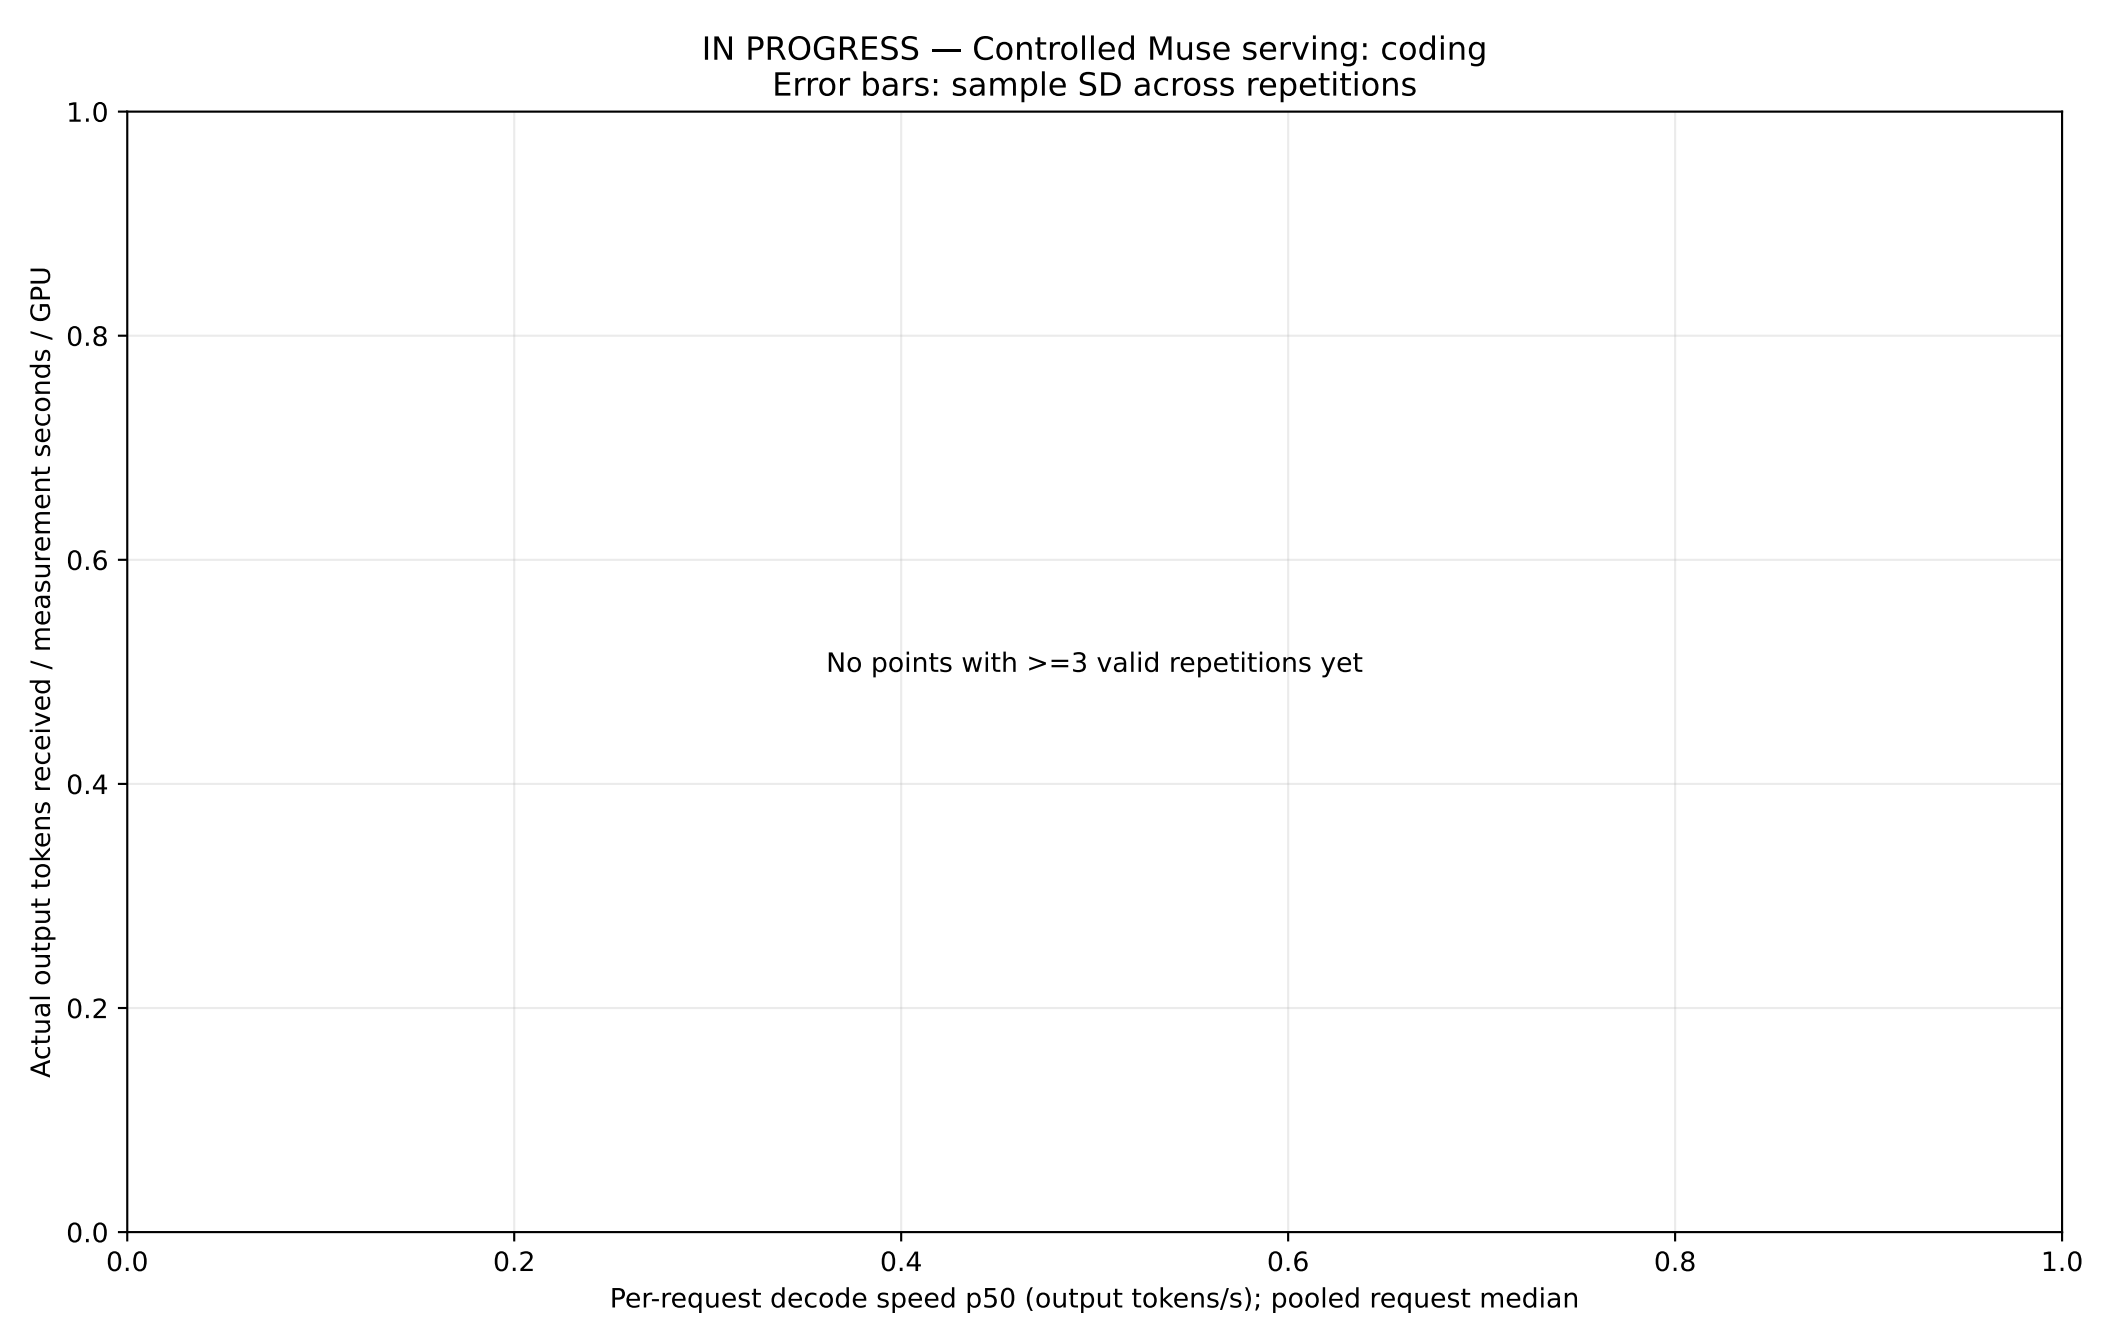
IN PROGRESS — Controlled Muse serving: coding
Error bars: sample SD across repetitions
No points with >=3 valid repetitions yet
Actual output tokens received / measurement seconds / GPU
Per-request decode speed p50 (output tokens/s); pooled request median
0.0
0.2
0.4
0.6
0.8
1.0