Muse Spark: Output Throughput vs. Interactivity

8,192 input / 1,024 output tokens | cache OFF | speculation OFF | FP8 KV

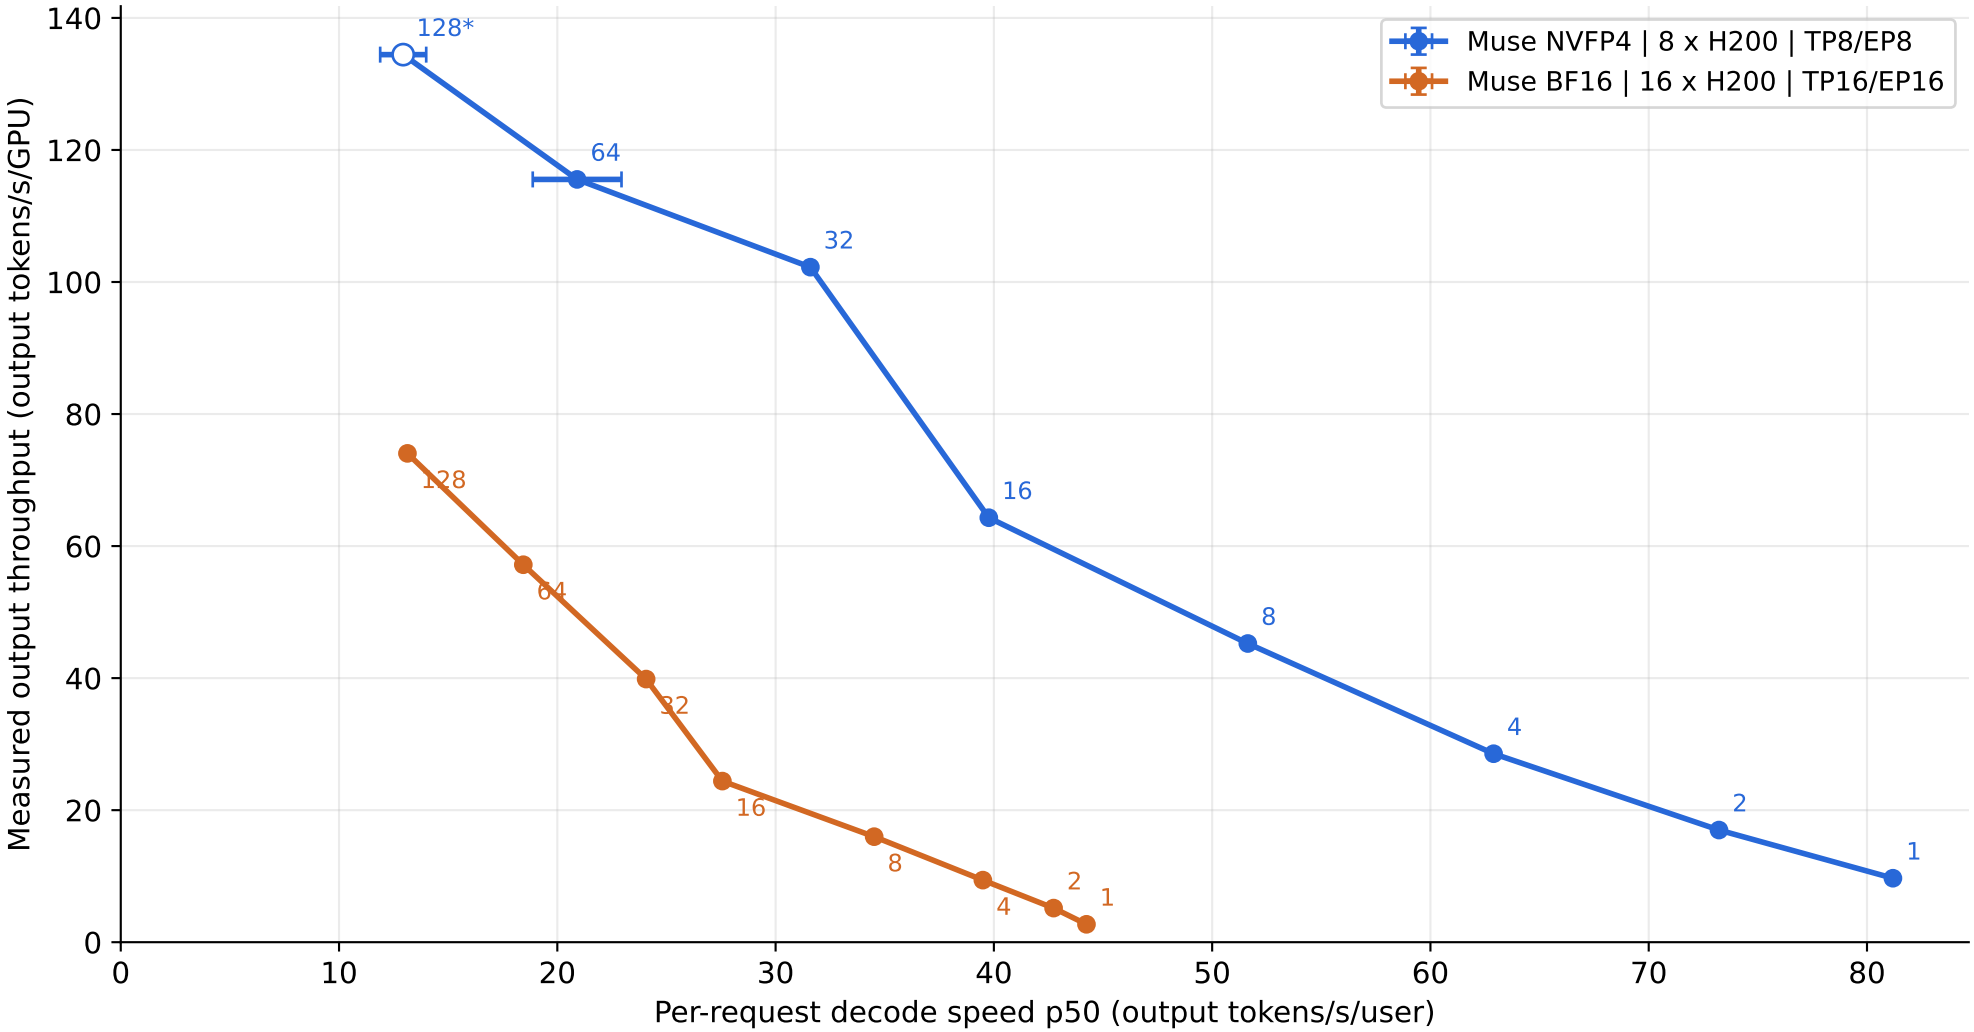

Labels: offered closed-loop concurrency | error bars: repetition SD | 3 repeats; NVFP4 c4/c8: 5

* Hollow point: TTFT p95 > 60 s or E2E p95 > 300 s. BF16 c256 omitted: only 1/3 valid repeats.

Serving windows include prefill + decode; BF16 prefill CUDA graphs disabled. Not a matched DeepInfer decode benchmark.