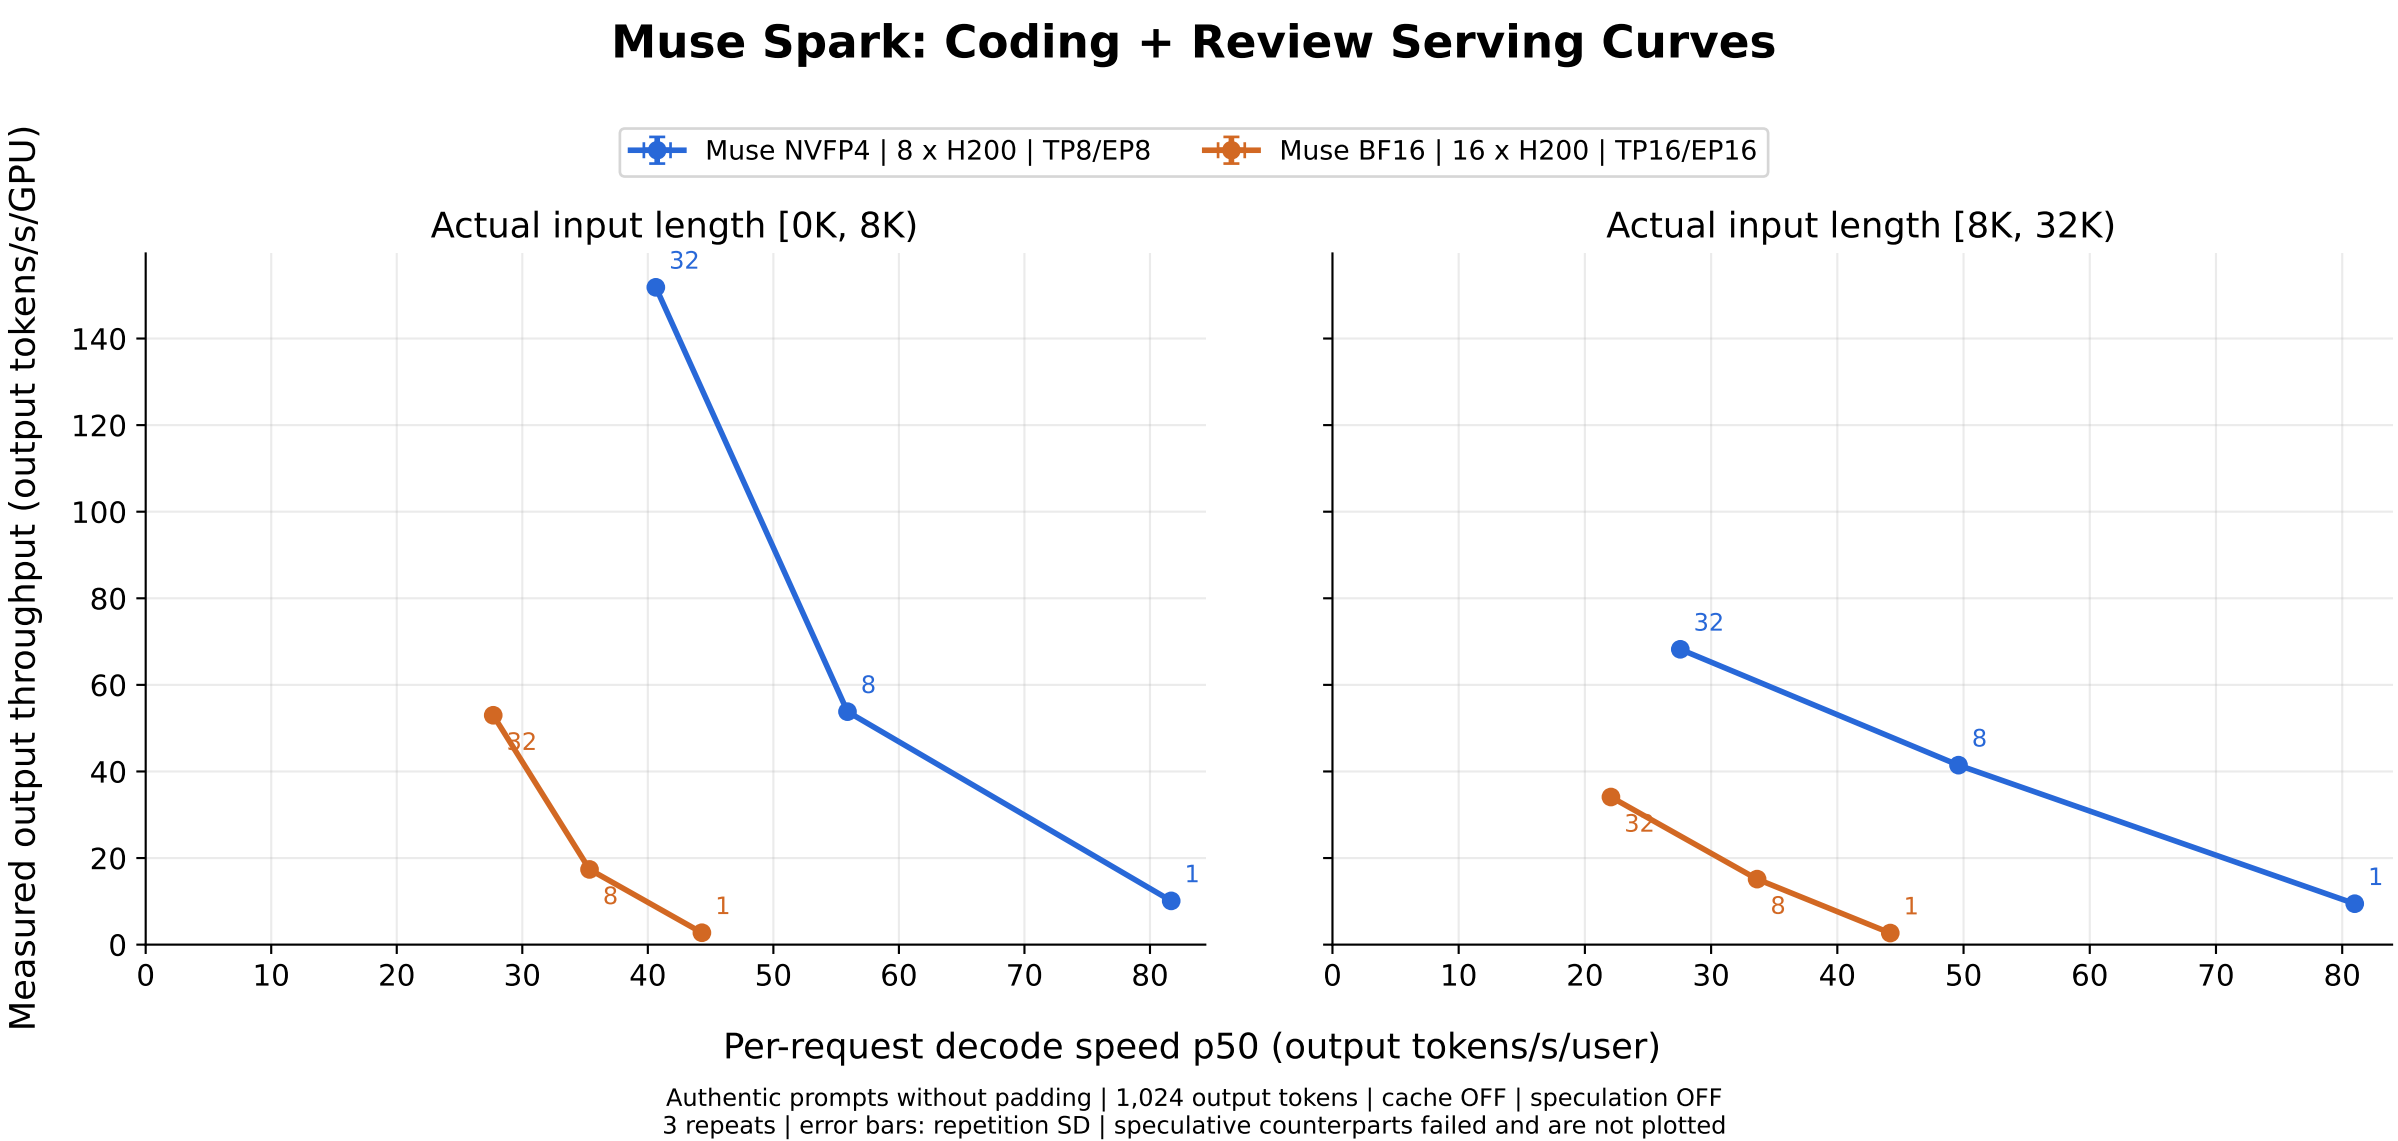
Muse Spark: Coding + Review Serving Curves
Muse NVFP4 | 8 x H200 | TP8/EP8
Muse BF16 | 16 x H200 | TP16/EP16
Actual input length [0K, 8K)
Actual input length [8K, 32K)
Measured output throughput (output tokens/s/GPU)
Per-request decode speed p50 (output tokens/s/user)
0
20
40
60
80
100
120
140
0
10
20
30
40
50
60
70
80
32
8
1
Authentic prompts without padding | 1,024 output tokens | cache OFF | speculation OFF
3 repeats | error bars: repetition SD | speculative counterparts failed and are not plotted