Muse Spark: Long-context Throughput vs. Interactivity

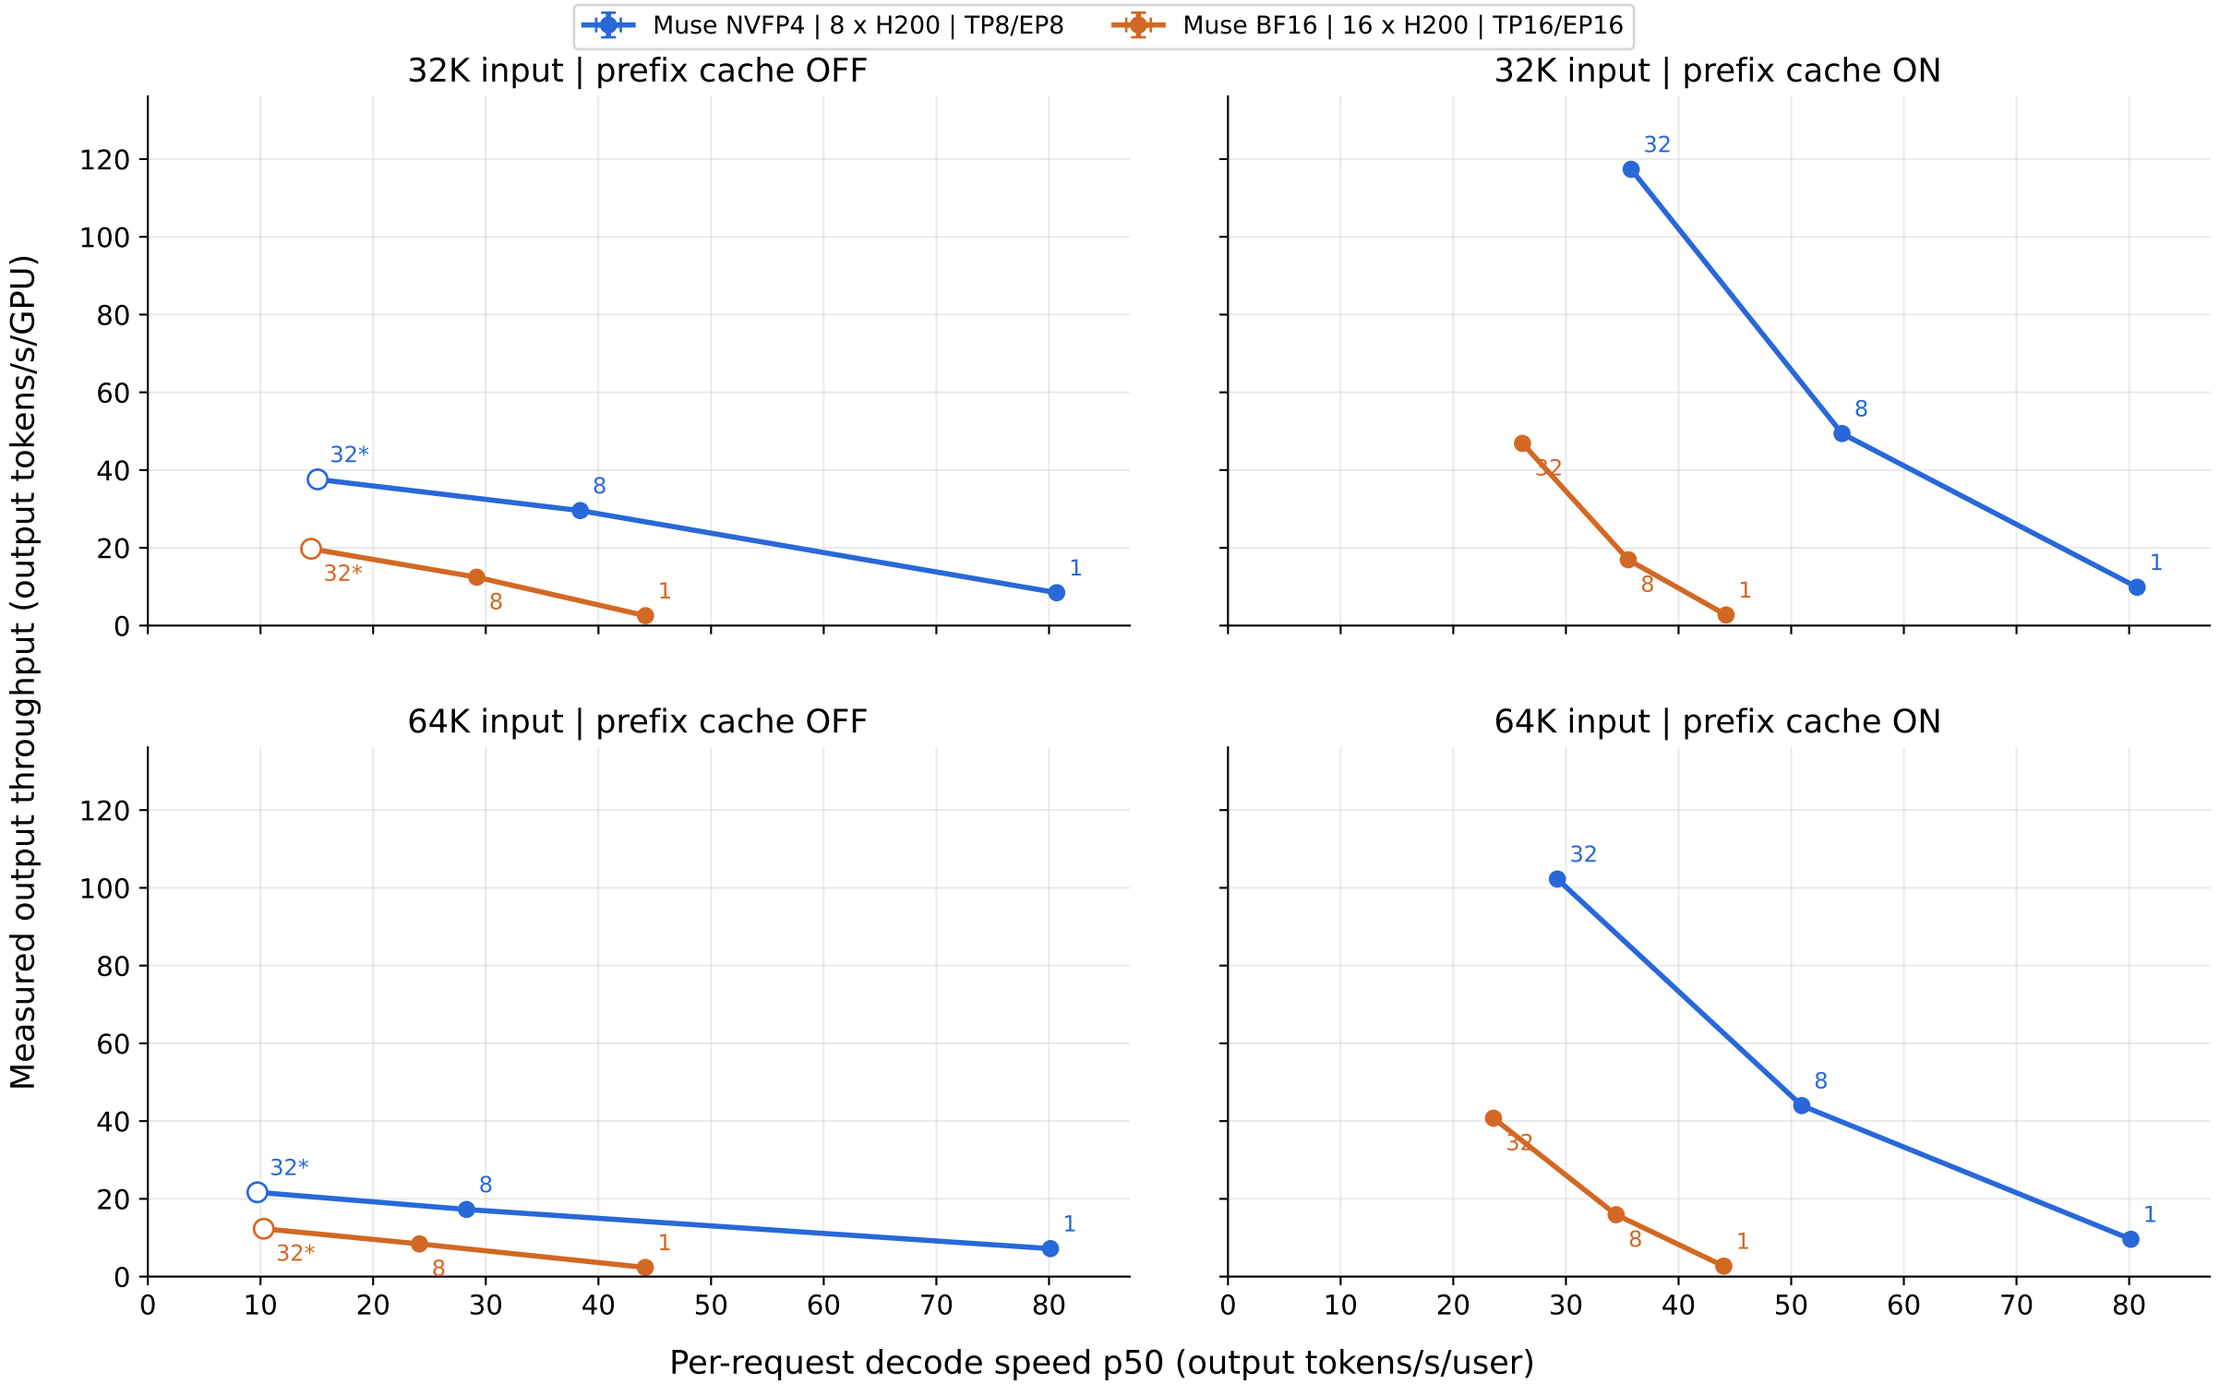

1,024 output tokens | speculation OFF | FP8 KV | labels: concurrency | 3 repeats | error bars: repetition SD
Cache ON: nominal 90% common prefix; measured hit rates in detailed_rounds.csv. * Hollow points exceed latency limits.