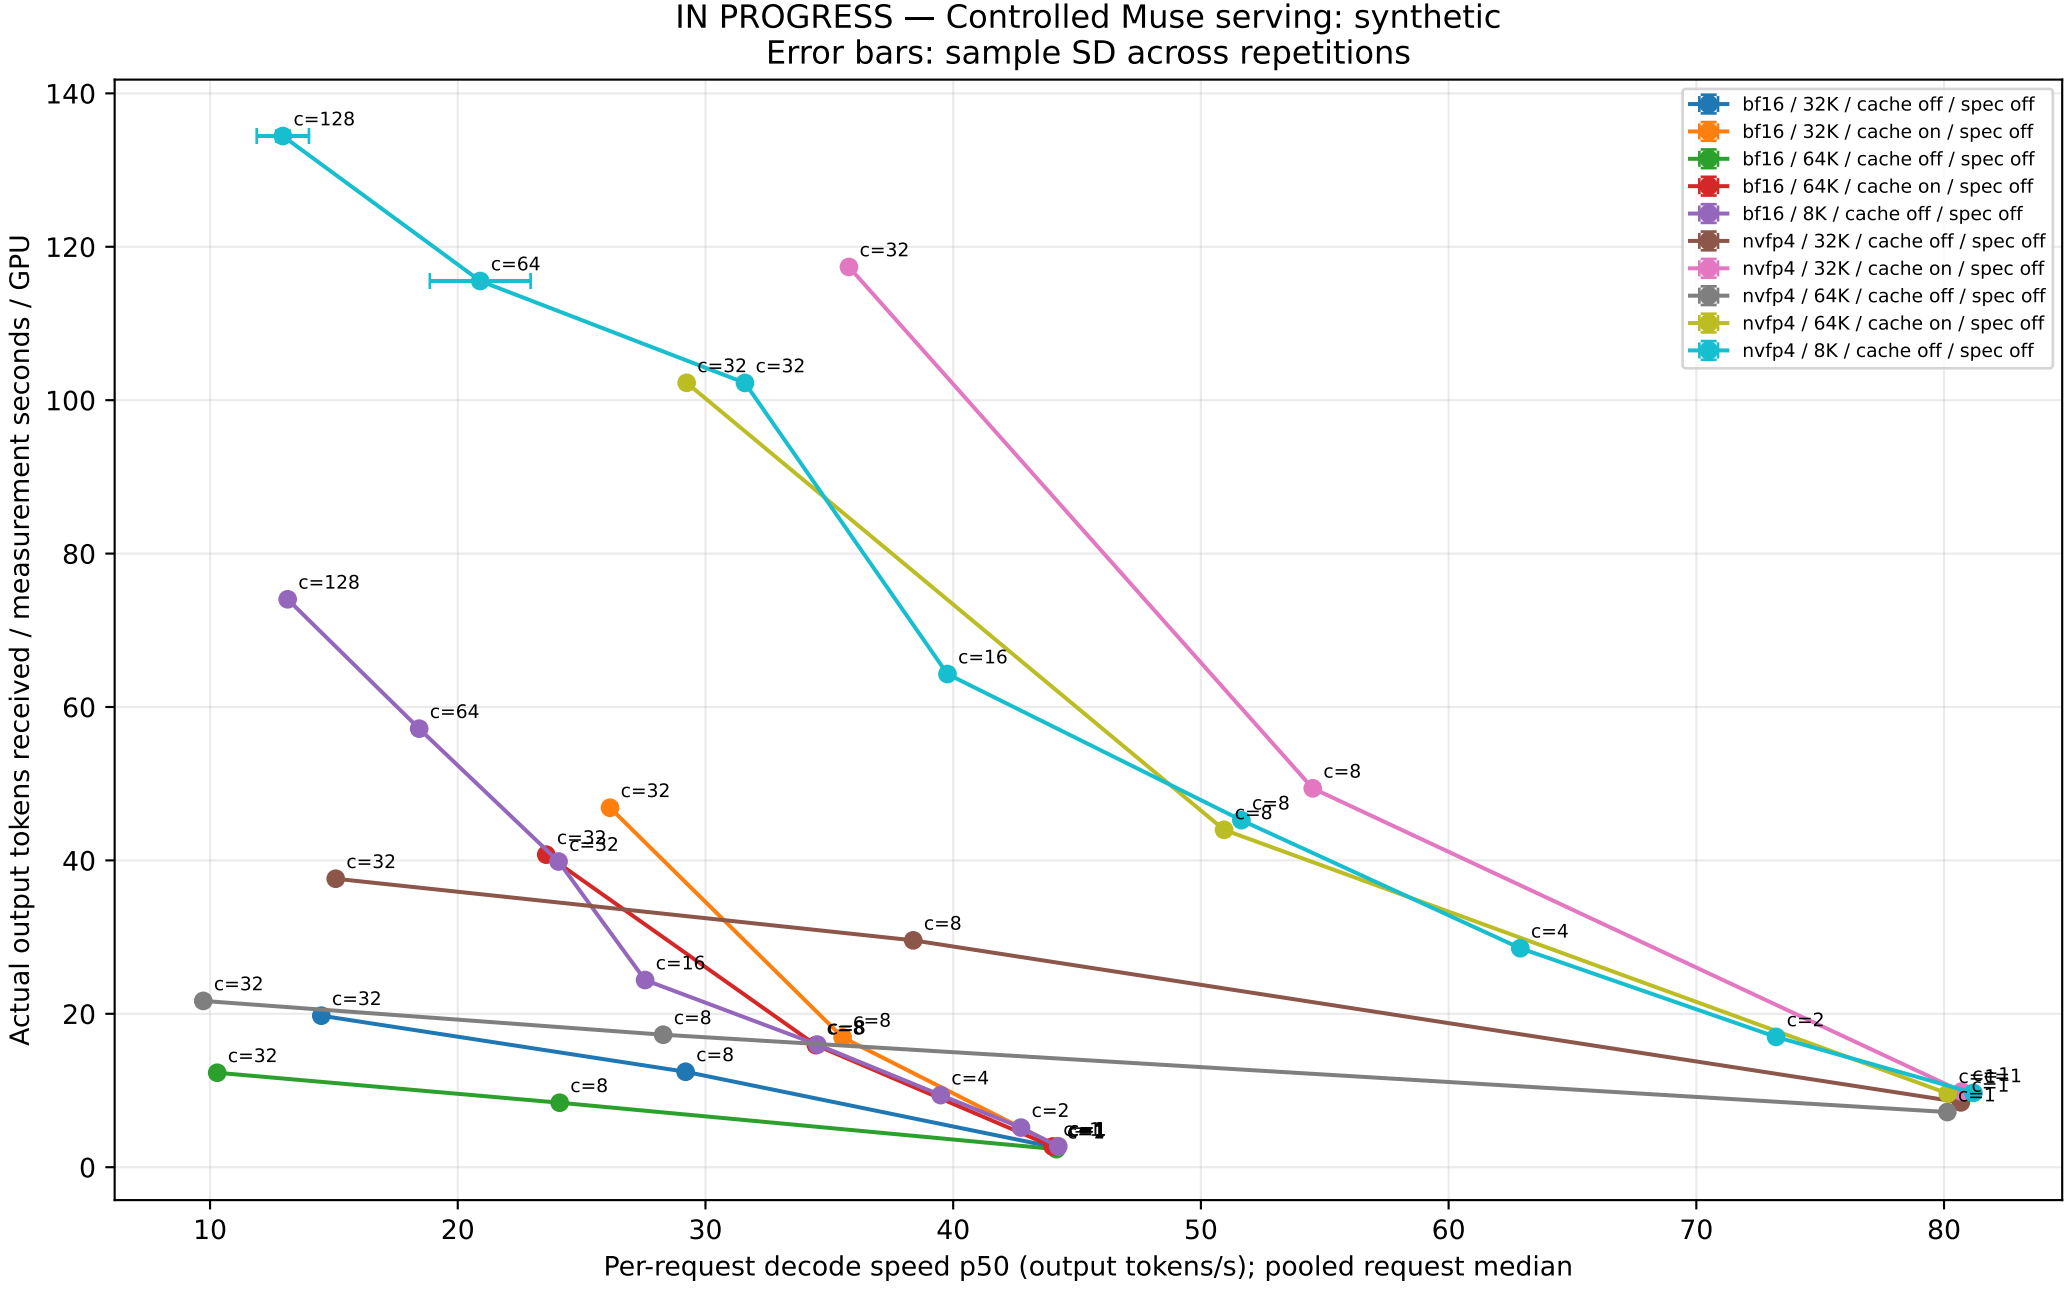
IN PROGRESS — Controlled Muse serving: synthetic
Error bars: sample SD across repetitions
Actual output tokens received / measurement seconds / GPU
Per-request decode speed p50 (output tokens/s); pooled request median
bf16 / 32K / cache off / spec off
bf16 / 32K / cache on / spec off
bf16 / 64K / cache off / spec off
bf16 / 64K / cache on / spec off
bf16 / 8K / cache off / spec off
nvfp4 / 32K / cache off / spec off
nvfp4 / 32K / cache on / spec off
nvfp4 / 64K / cache off / spec off
nvfp4 / 64K / cache on / spec off
nvfp4 / 8K / cache off / spec off