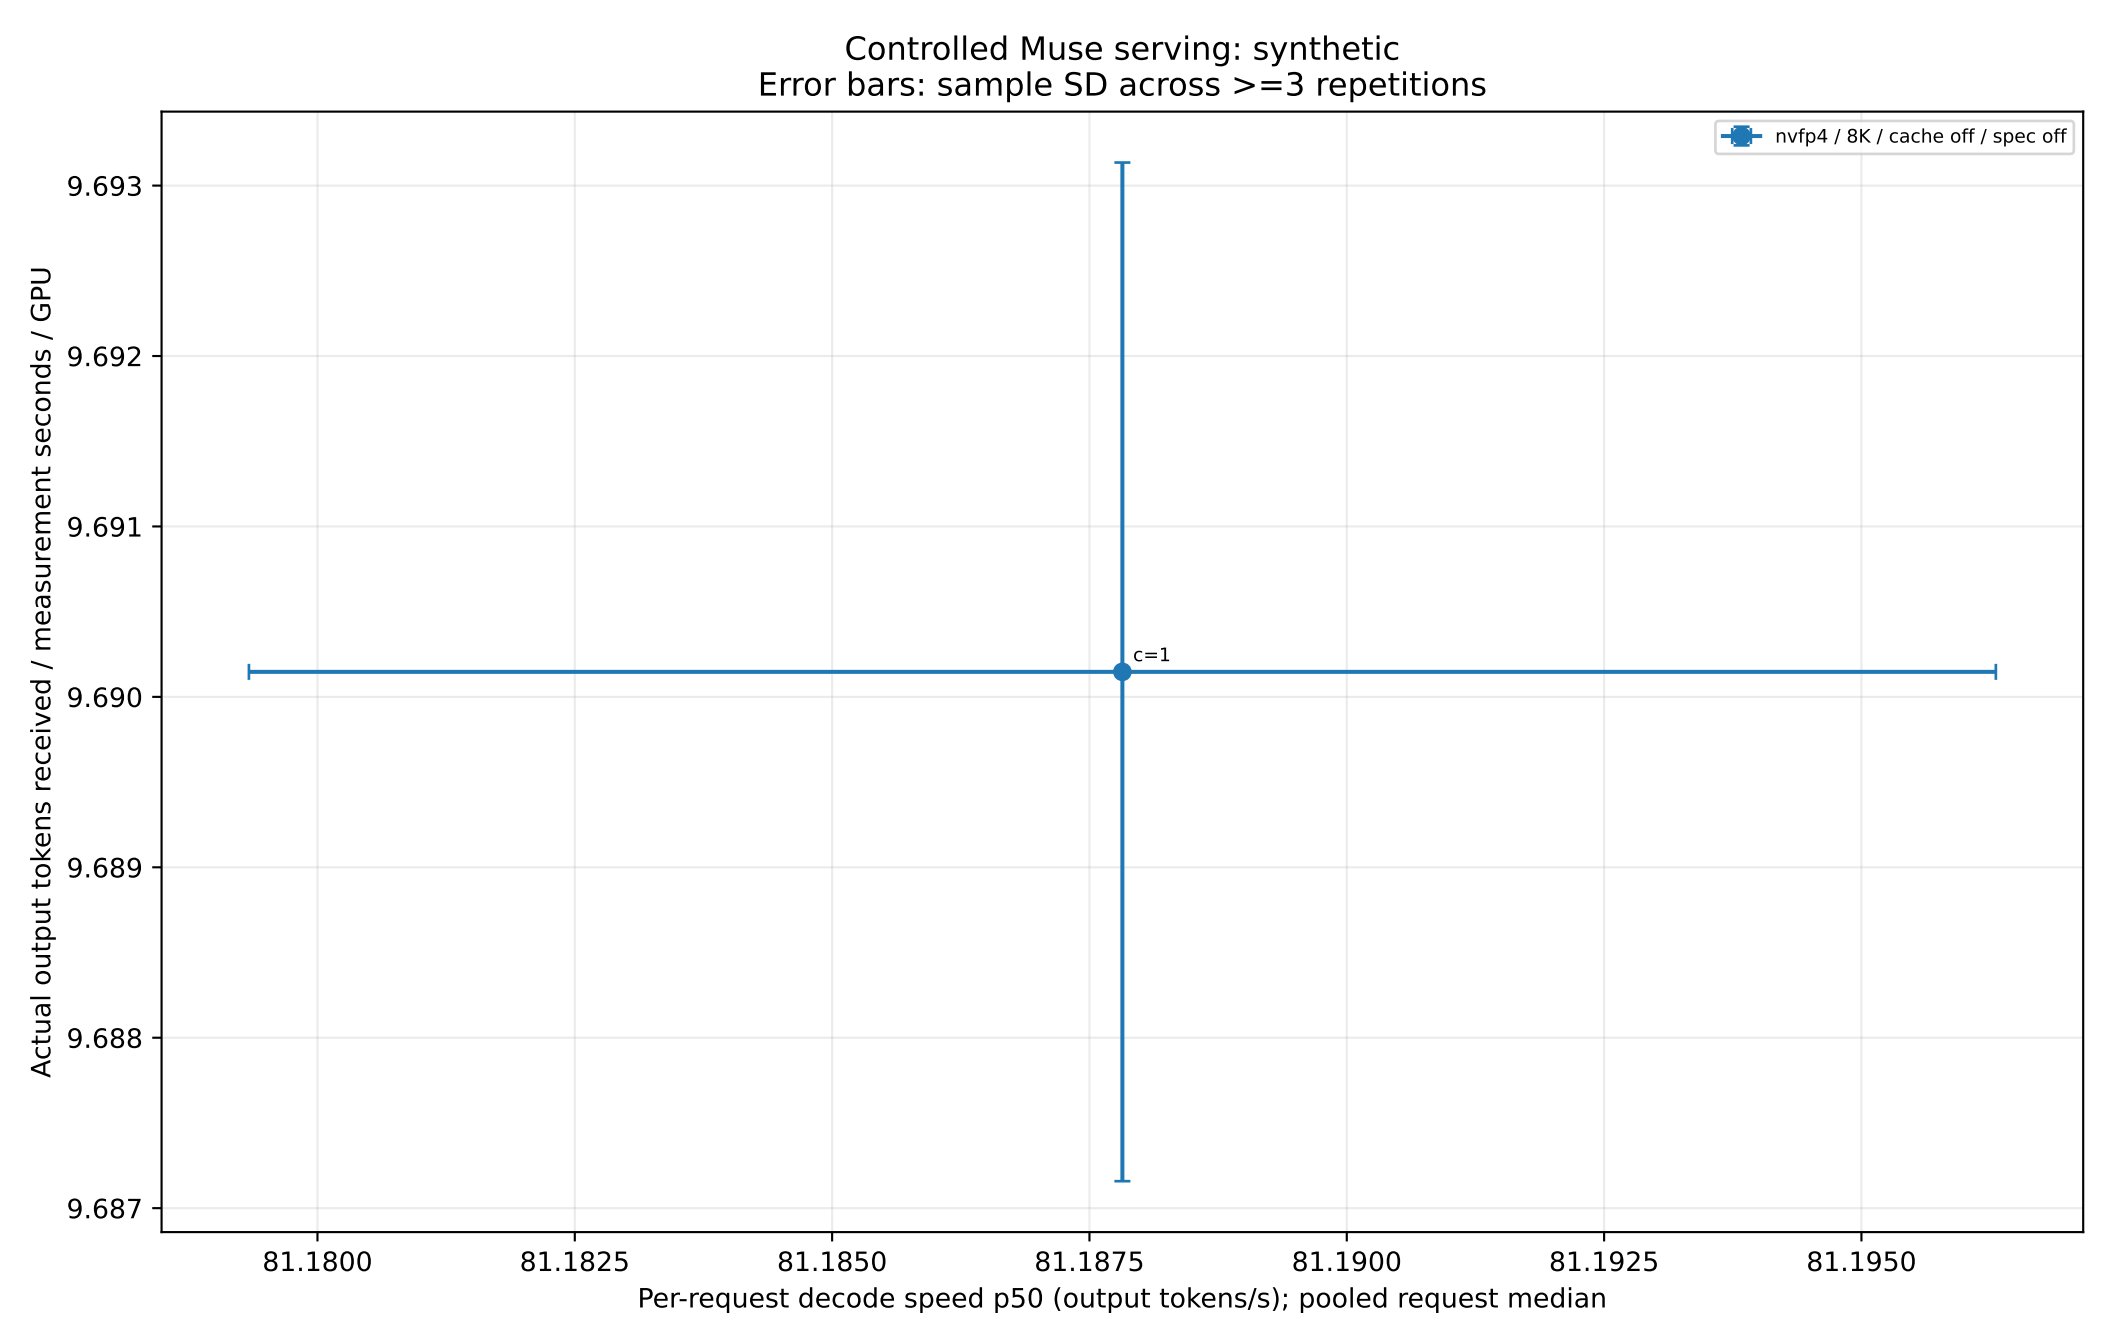
Controlled Muse serving: synthetic
Error bars: sample SD across >=3 repetitions
nvfp4 / 8K / cache off / spec off
c=1
Actual output tokens received / measurement seconds / GPU
9.693
9.692
9.691
9.690
9.689
9.688
9.687
81.1800
81.1825
81.1850
81.1875
81.1900
81.1925
81.1950
Per-request decode speed p50 (output tokens/s); pooled request median