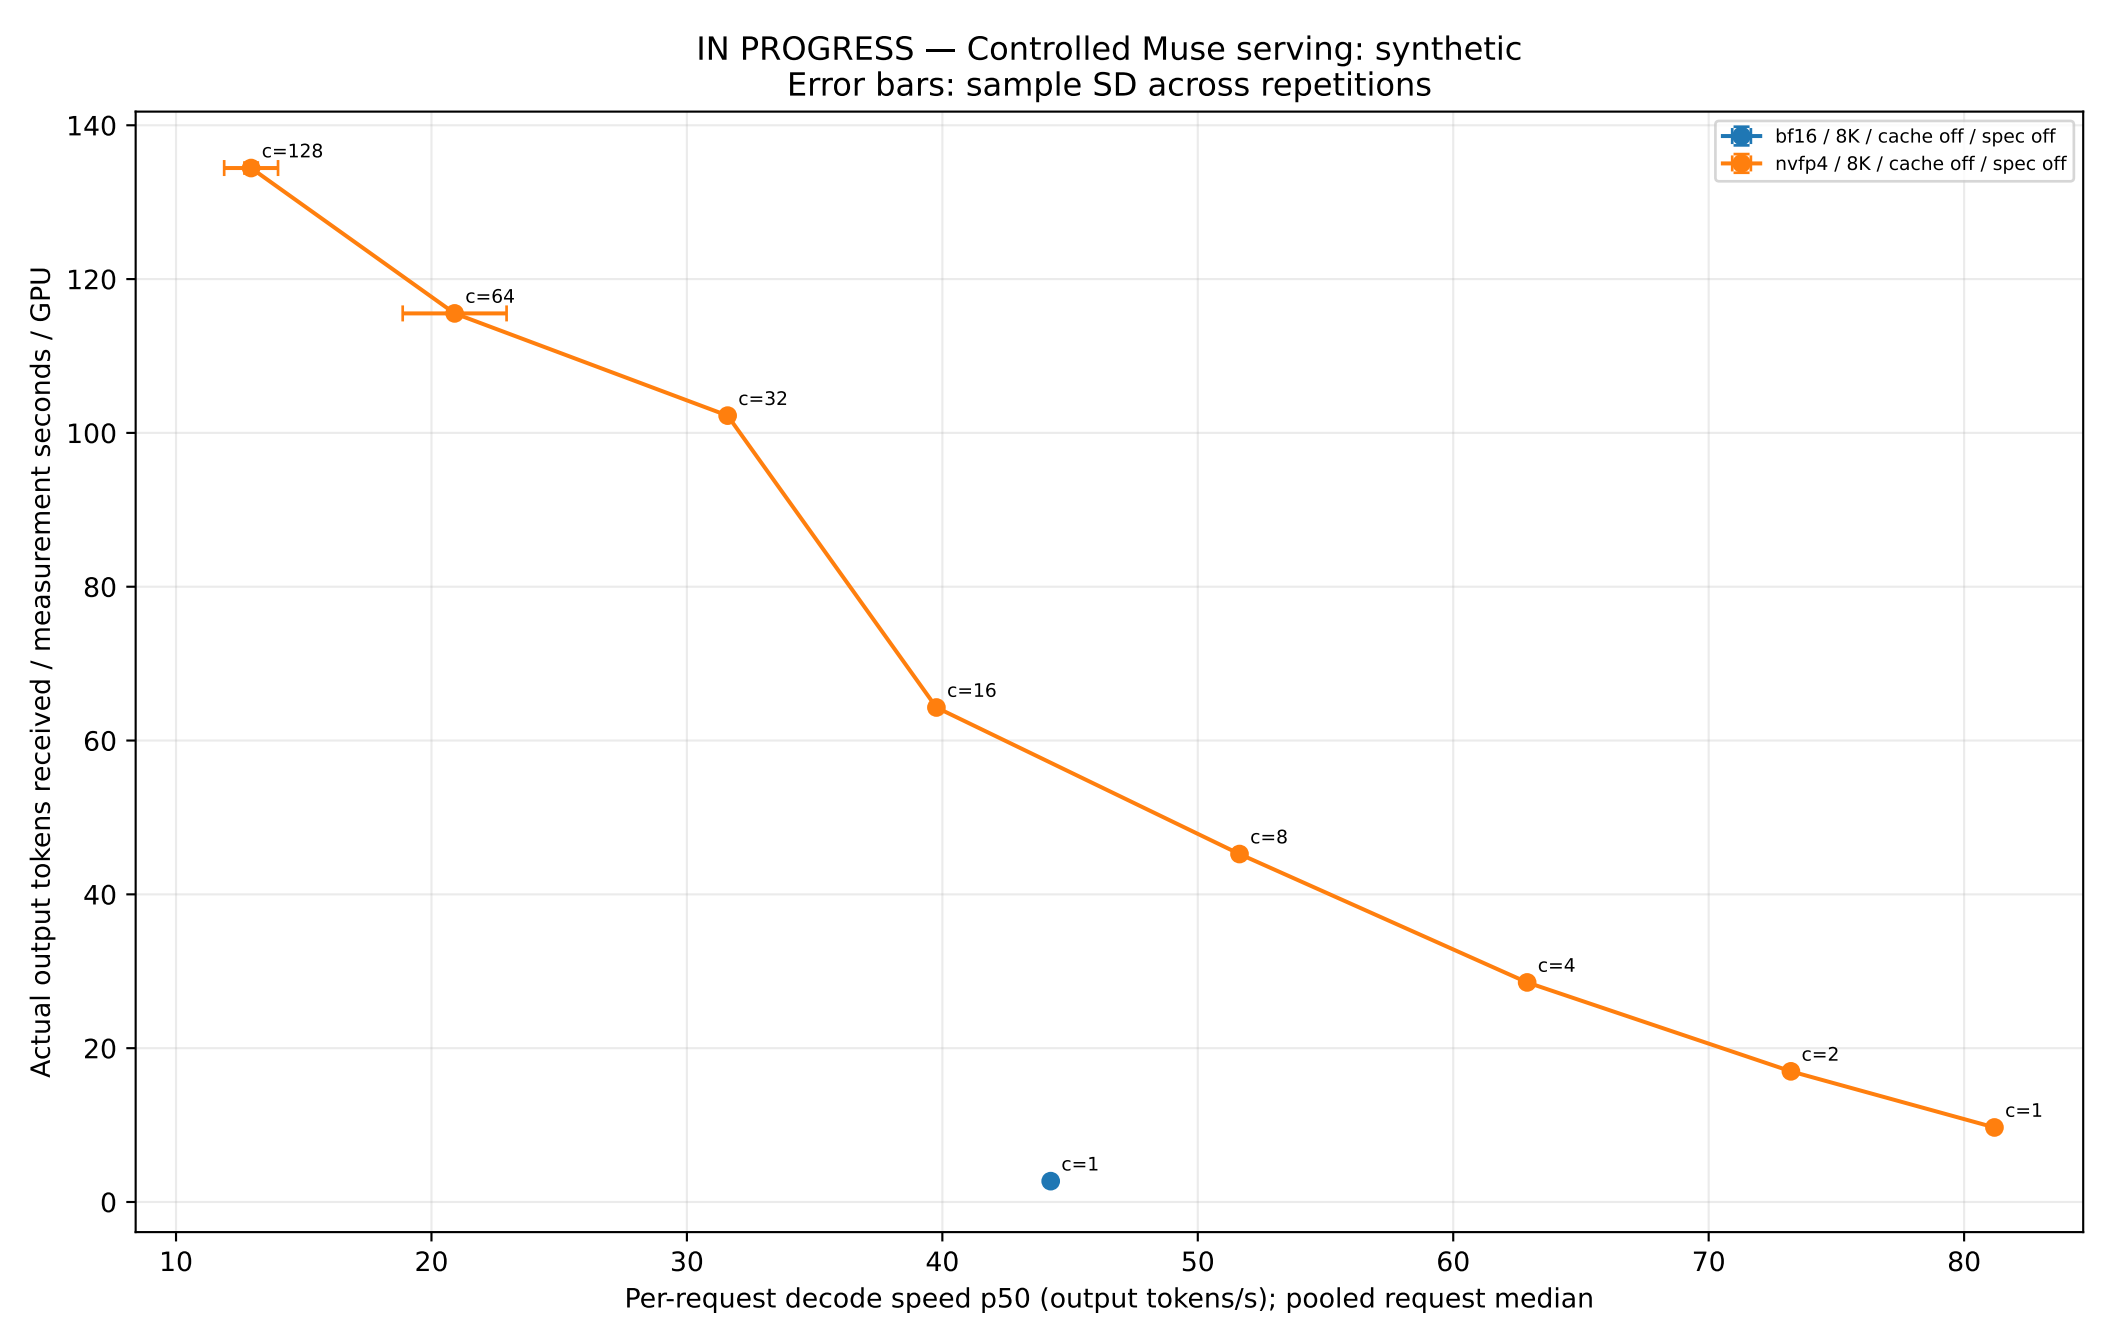
IN PROGRESS — Controlled Muse serving: synthetic
Error bars: sample SD across repetitions
bf16 / 8K / cache off / spec off
nvfp4 / 8K / cache off / spec off
c=128
c=64
c=32
c=16
c=8
c=4
c=2
c=1
c=1
Actual output tokens received / measurement seconds / GPU
Per-request decode speed p50 (output tokens/s); pooled request median
0
20
40
60
80
100
120
140
10
20
30
40
50
60
70
80