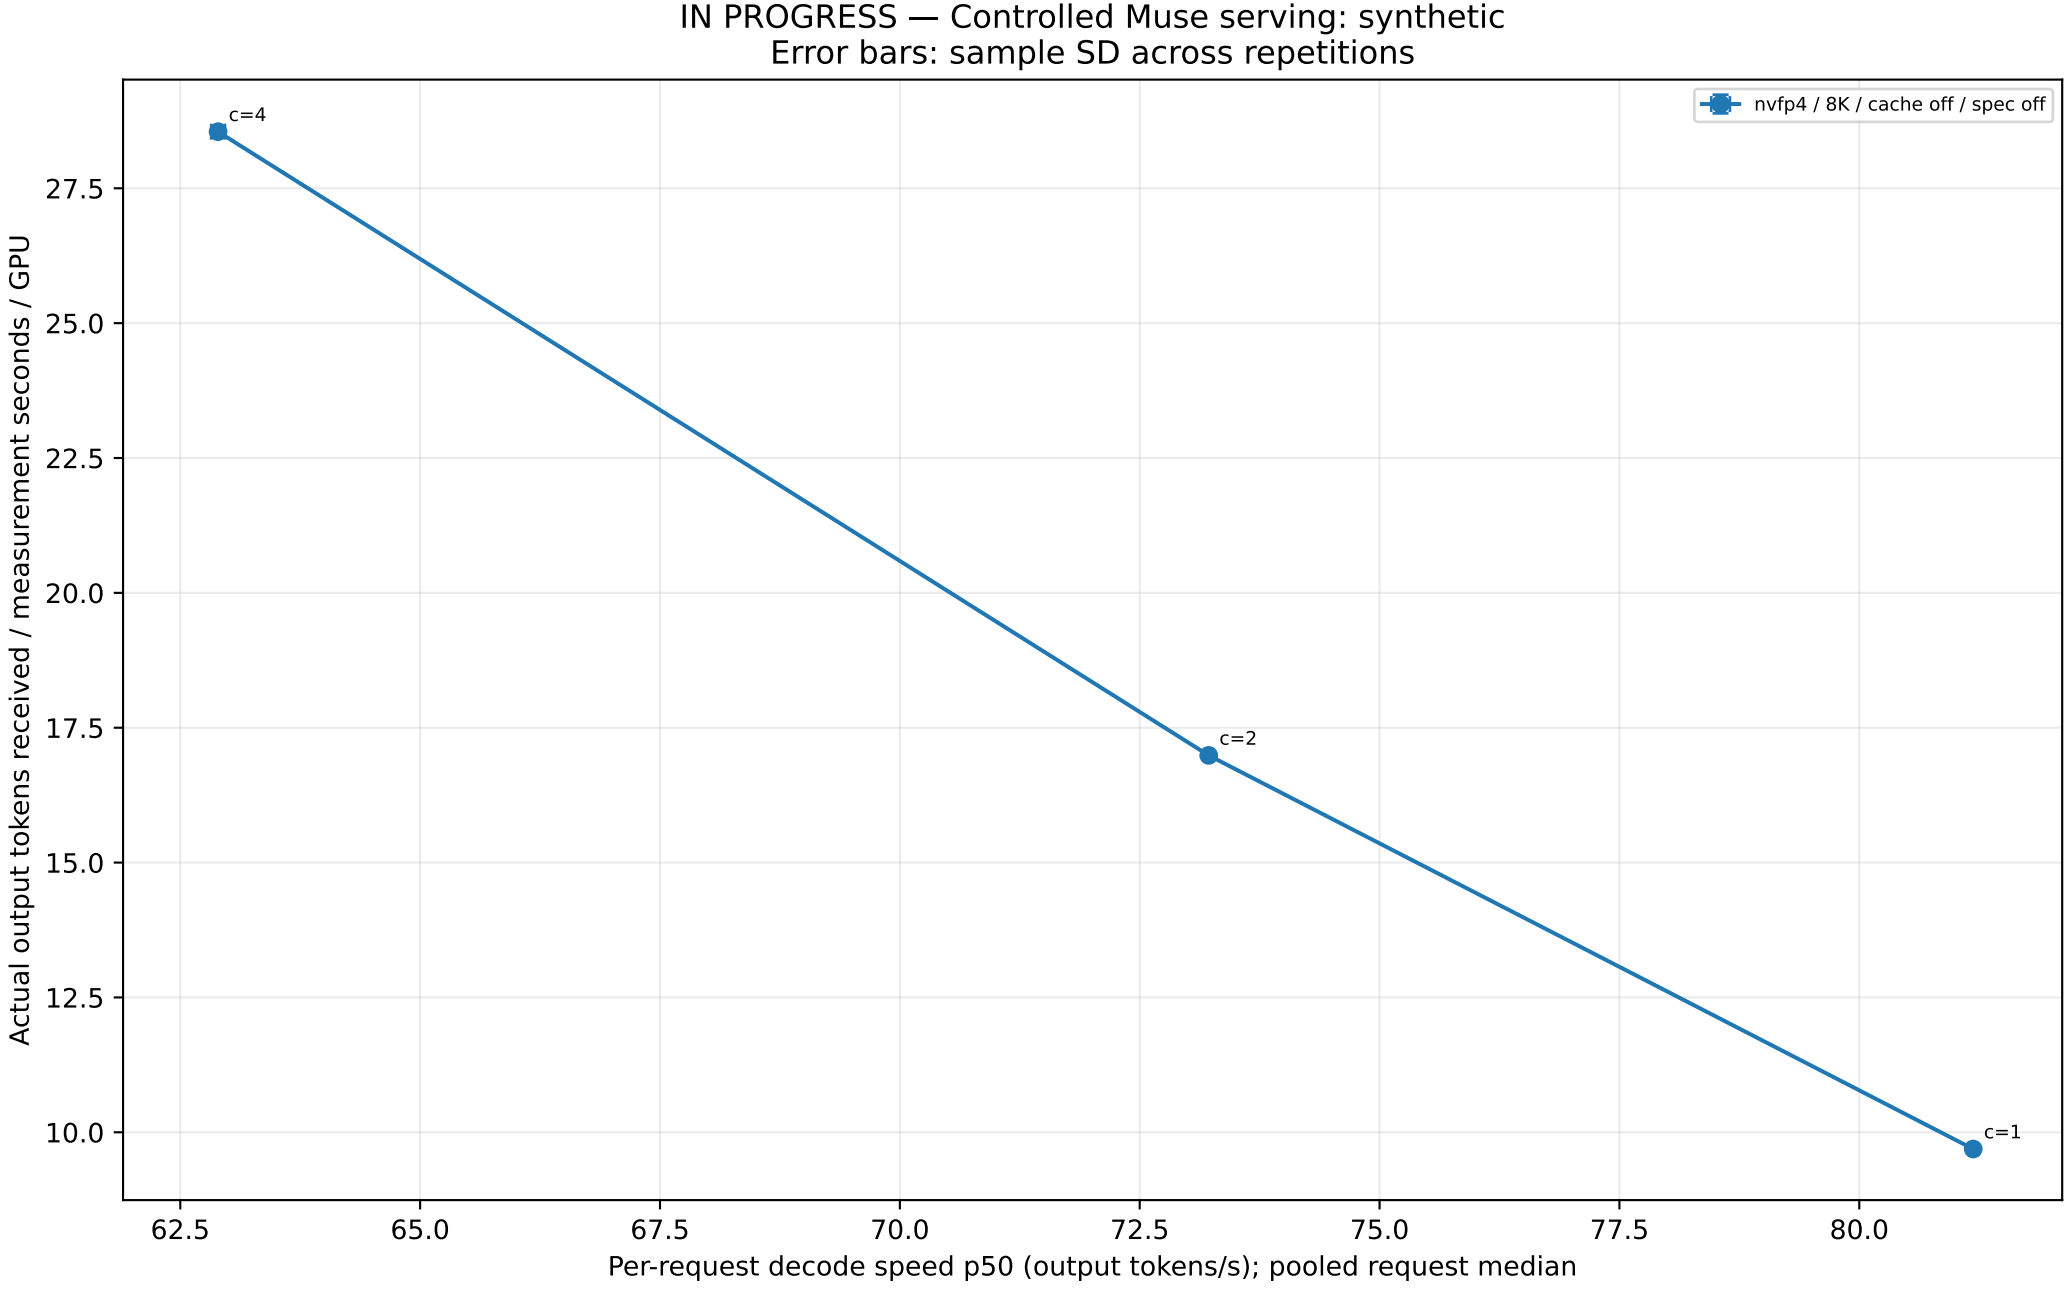
IN PROGRESS — Controlled Muse serving: synthetic
Error bars: sample SD across repetitions
nvfp4 / 8K / cache off / spec off
c=4
c=2
c=1
Actual output tokens received / measurement seconds / GPU
Per-request decode speed p50 (output tokens/s); pooled request median
10.0
12.5
15.0
17.5
20.0
22.5
25.0
27.5
62.5
65.0
67.5
70.0
72.5
75.0
77.5
80.0