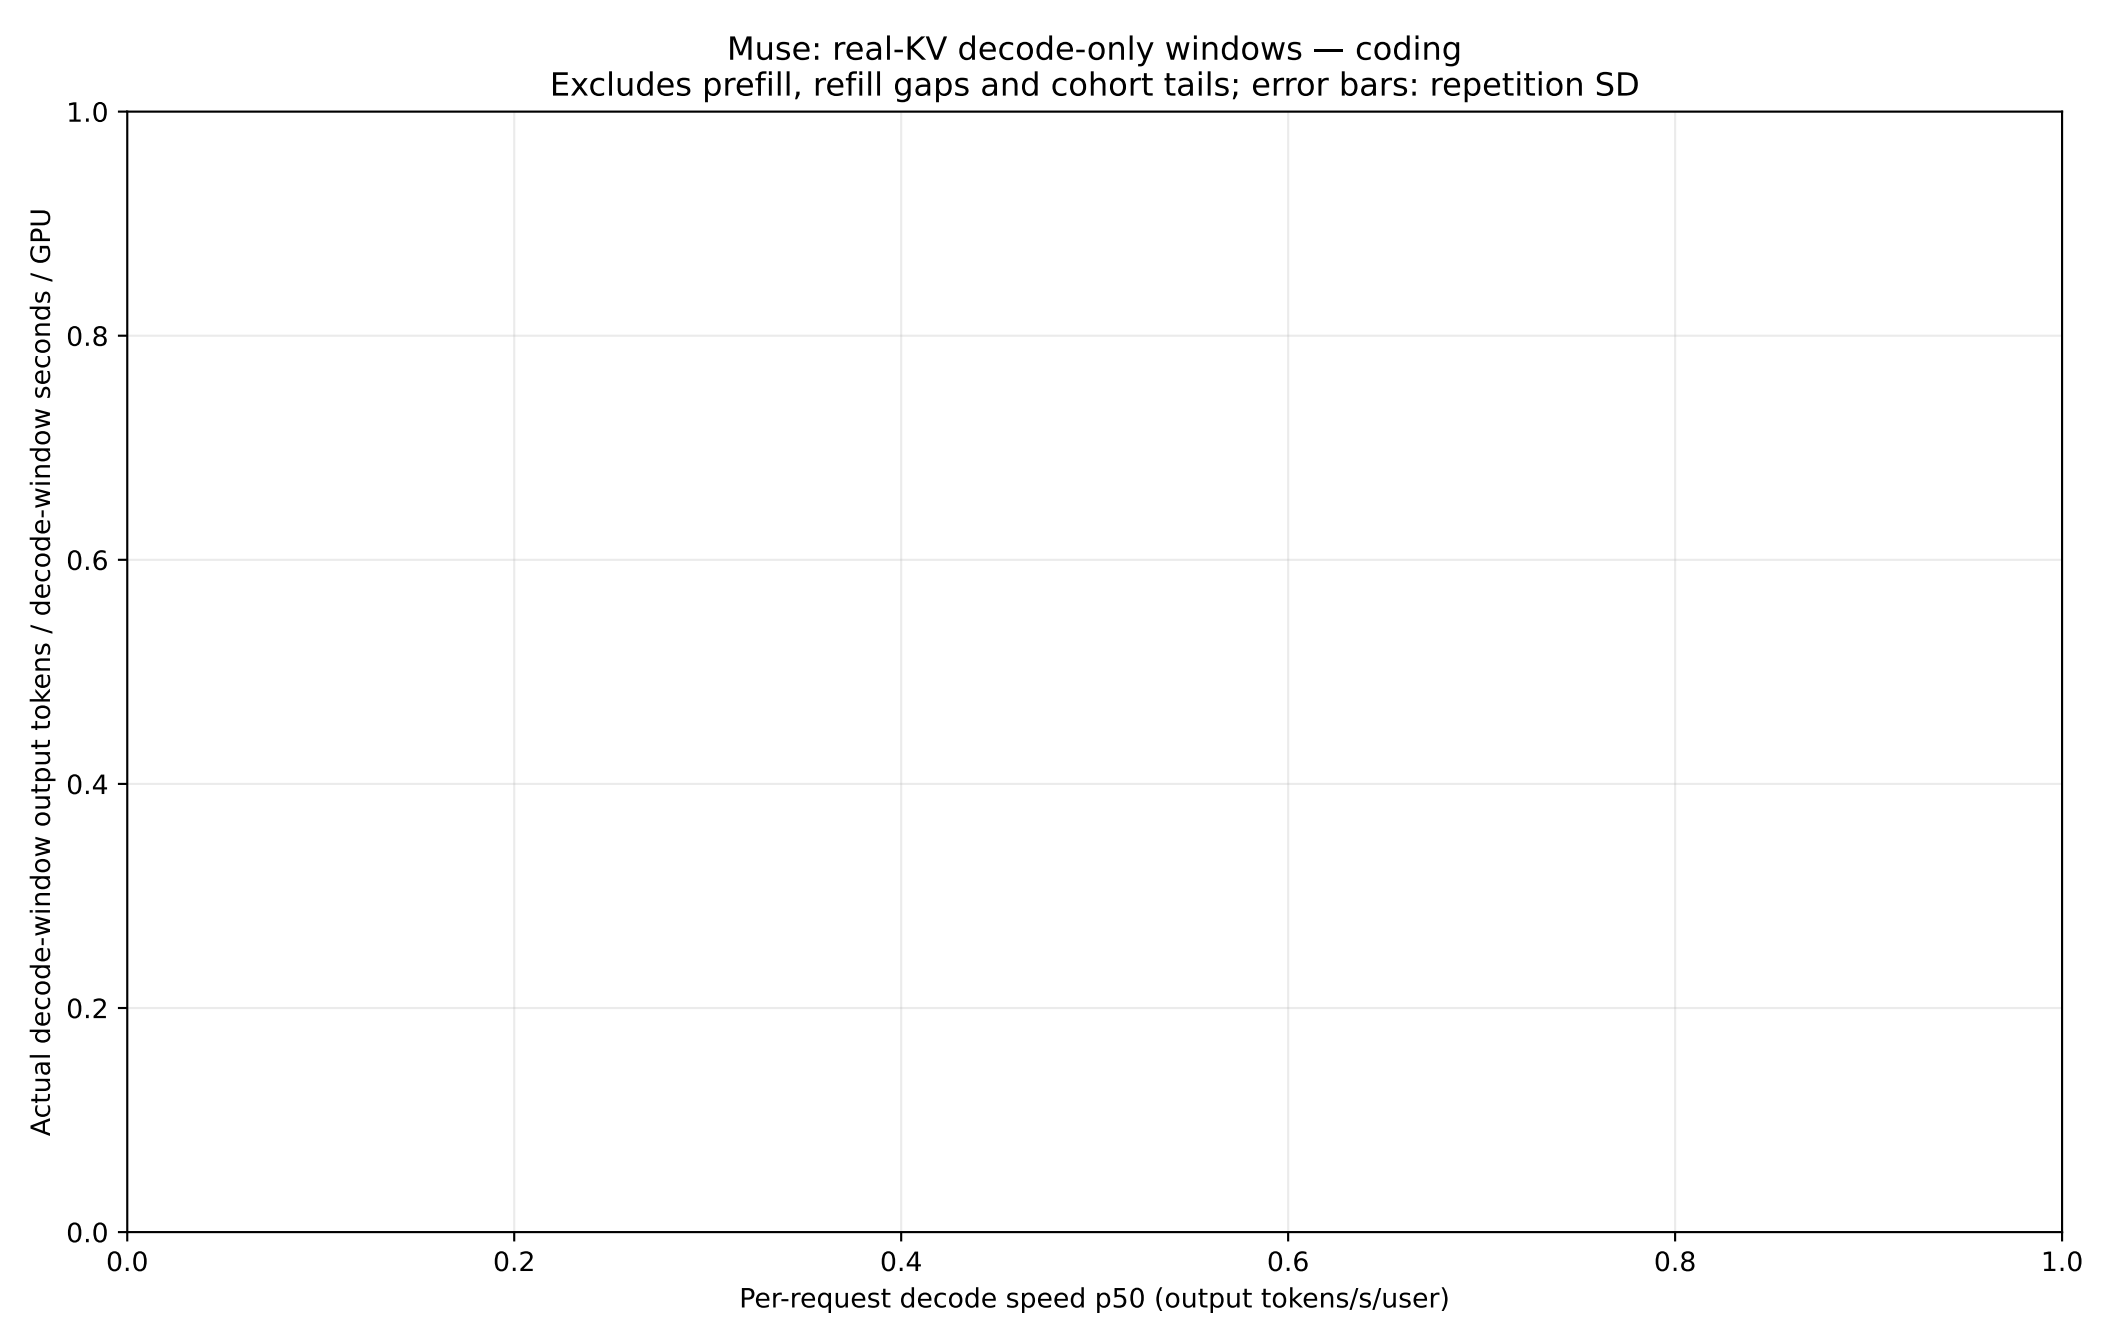
Muse: real-KV decode-only windows — coding
Excludes prefill, refill gaps and cohort tails; error bars: repetition SD
Actual decode-window output tokens / decode-window seconds / GPU
Per-request decode speed p50 (output tokens/s/user)
0.0
0.2
0.4
0.6
0.8
1.0