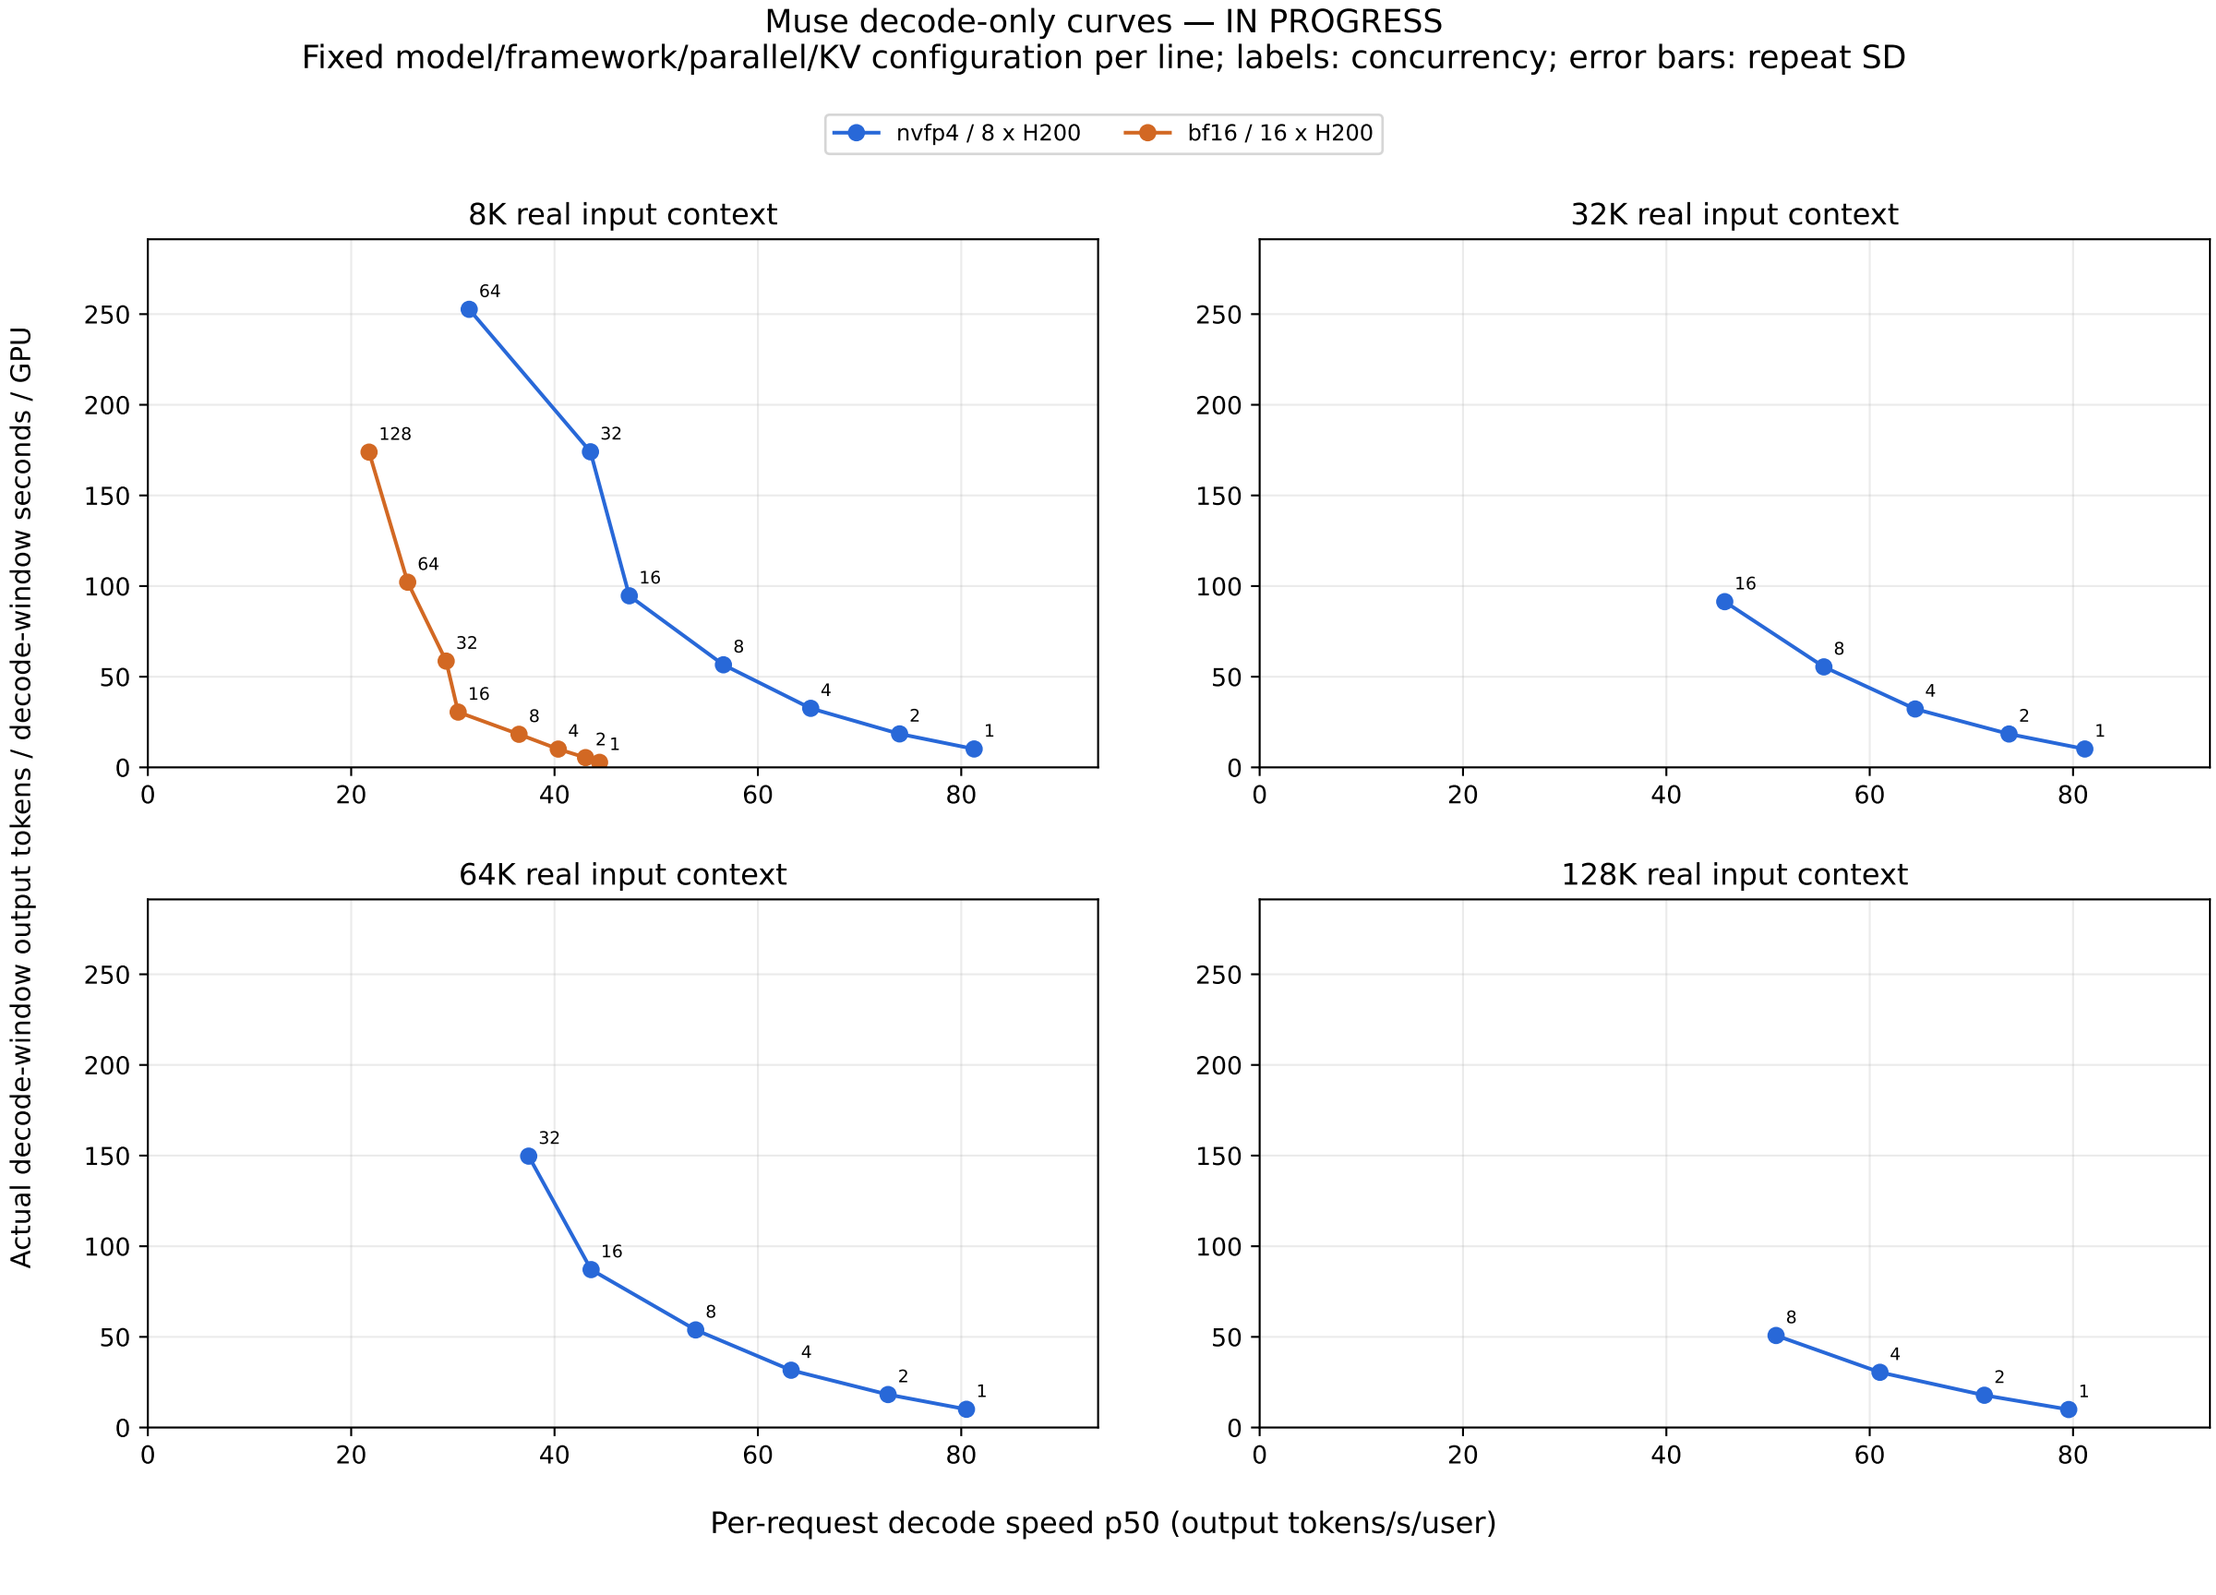
Muse decode-only curves — IN PROGRESS
Fixed model/framework/parallel/KV configuration per line; labels: concurrency; error bars: repeat SD
nvfp4 / 8 x H200
bf16 / 16 x H200
8K real input context
32K real input context
64K real input context
128K real input context
Actual decode-window output tokens / decode-window seconds / GPU
Per-request decode speed p50 (output tokens/s/user)
0
20
40
60
80
0
50
100
150
200
250
64
32
16
8
4
2
1
128
64
32
16
8
4
2
1
16
8
4
2
1
32
16
8
4
2
1
8
4
2
1