Muse and DeepInfer published decode curves — EXTERNAL REFERENCE ONLY

No hardware conversion, no inferred speedup ratio. Logarithmic Y axis.

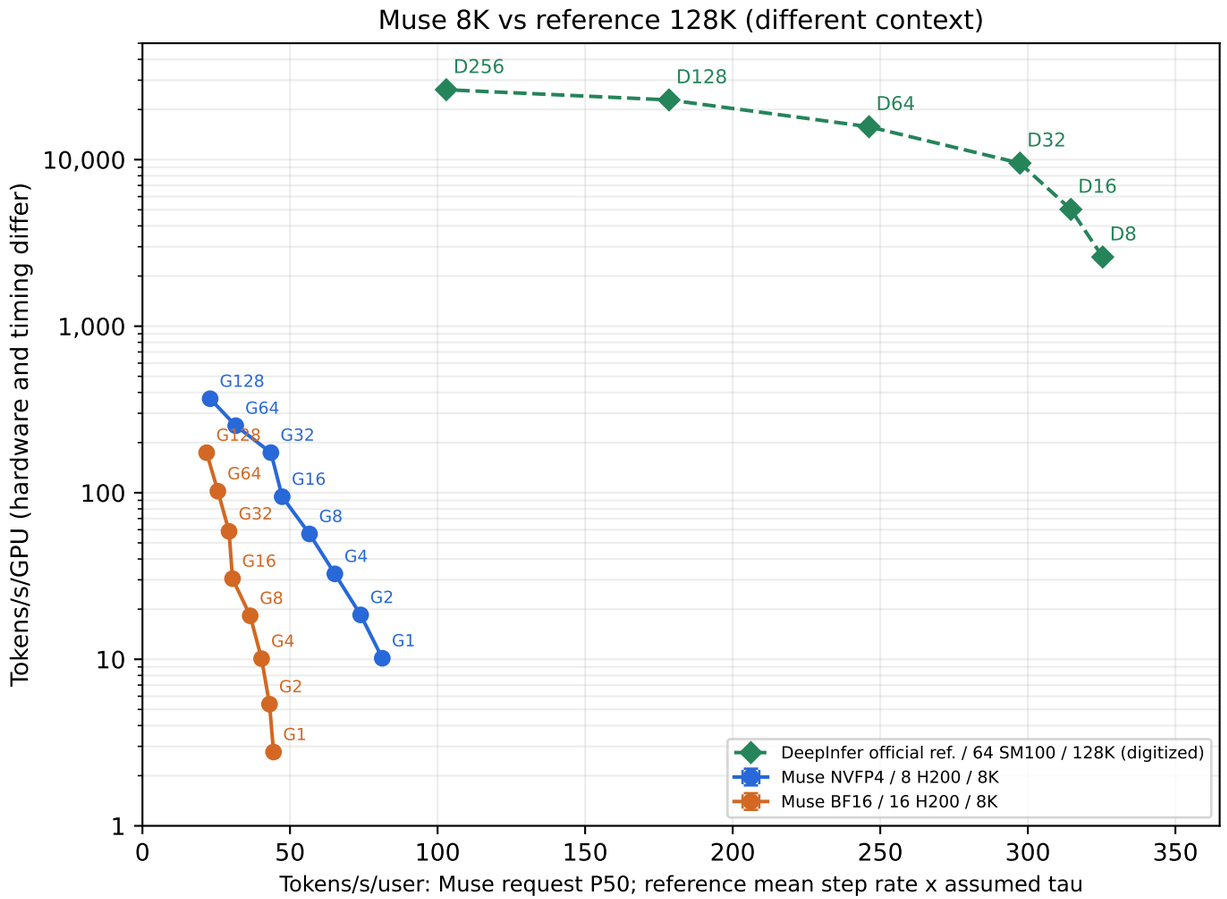

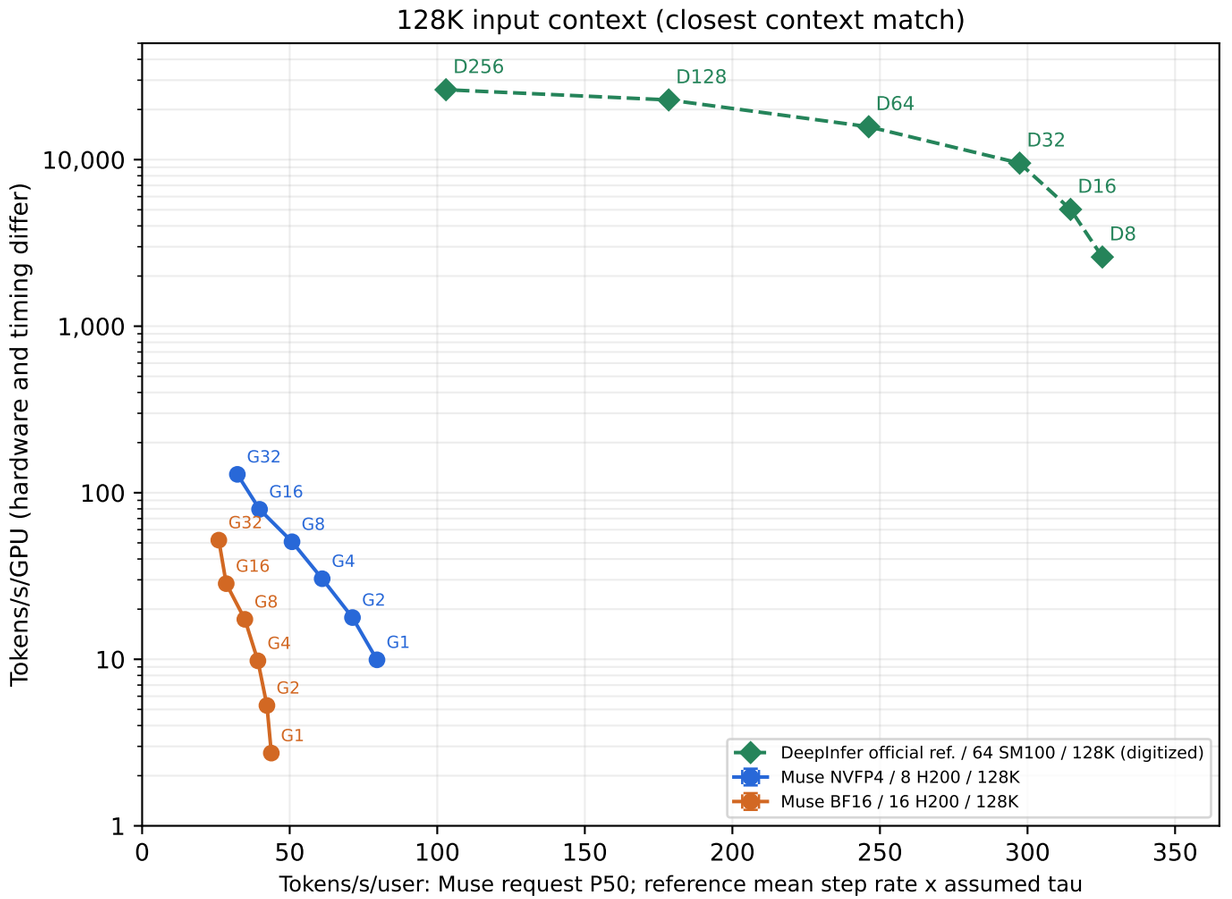

Muse: SGLang; real KV and streamed token increments; global concurrency G; 8/16 H200; speculative OFF; repeat SD error bars.
DeepInfer: DeepSeek V4.1 Flash; 64 SM100 in one scale-up domain; EP64/DP64; context 131,072; fixed nextn=3.
Reference labels D are requests PER DP RANK: D8 = 512 global, D256 = 16,384 global; not Muse G8/G256.
Reference axes use mean steps/s x tau, tau=1+3 x assumed 75%=3.25; not request P50 or native accepted-token throughput.
Reference values are digitized from the pinned official PNG; no raw repetitions or measurement error bars are available.

Source: deepseek-ai/DeepInfer @ 553741557a57; original image, Apache-2.0 license and pixel provenance included.