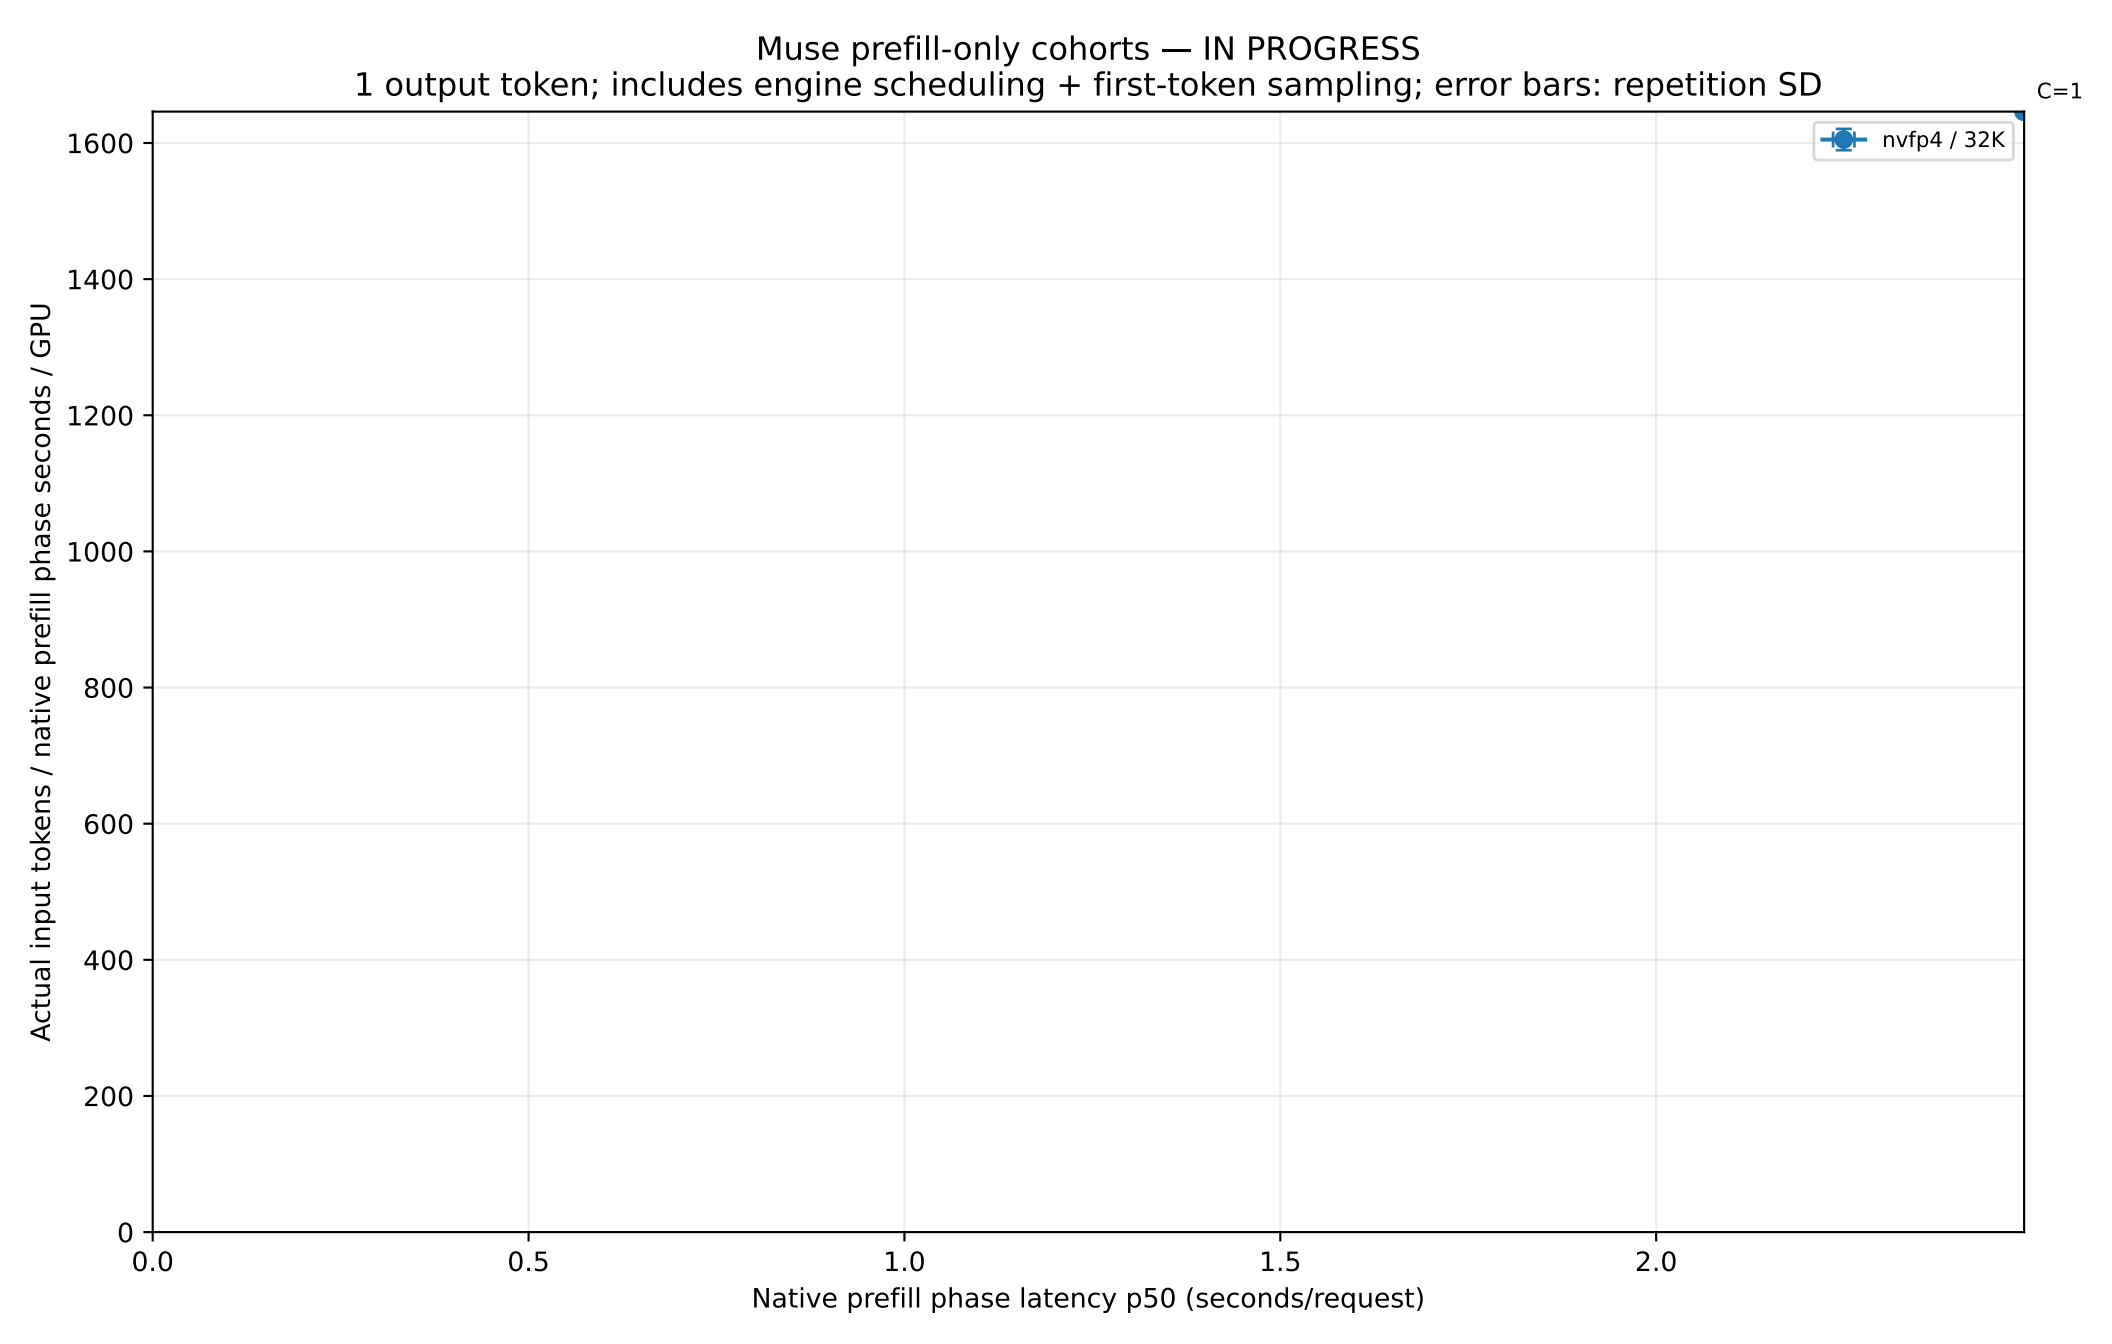
Muse prefill-only cohorts — IN PROGRESS
1 output token; includes engine scheduling + first-token sampling; error bars: repetition SD
C=1
nvfp4 / 32K
Actual input tokens / native prefill phase seconds / GPU
0
200
400
600
800
1000
1200
1400
1600
0.0
0.5
1.0
1.5
2.0
Native prefill phase latency p50 (seconds/request)