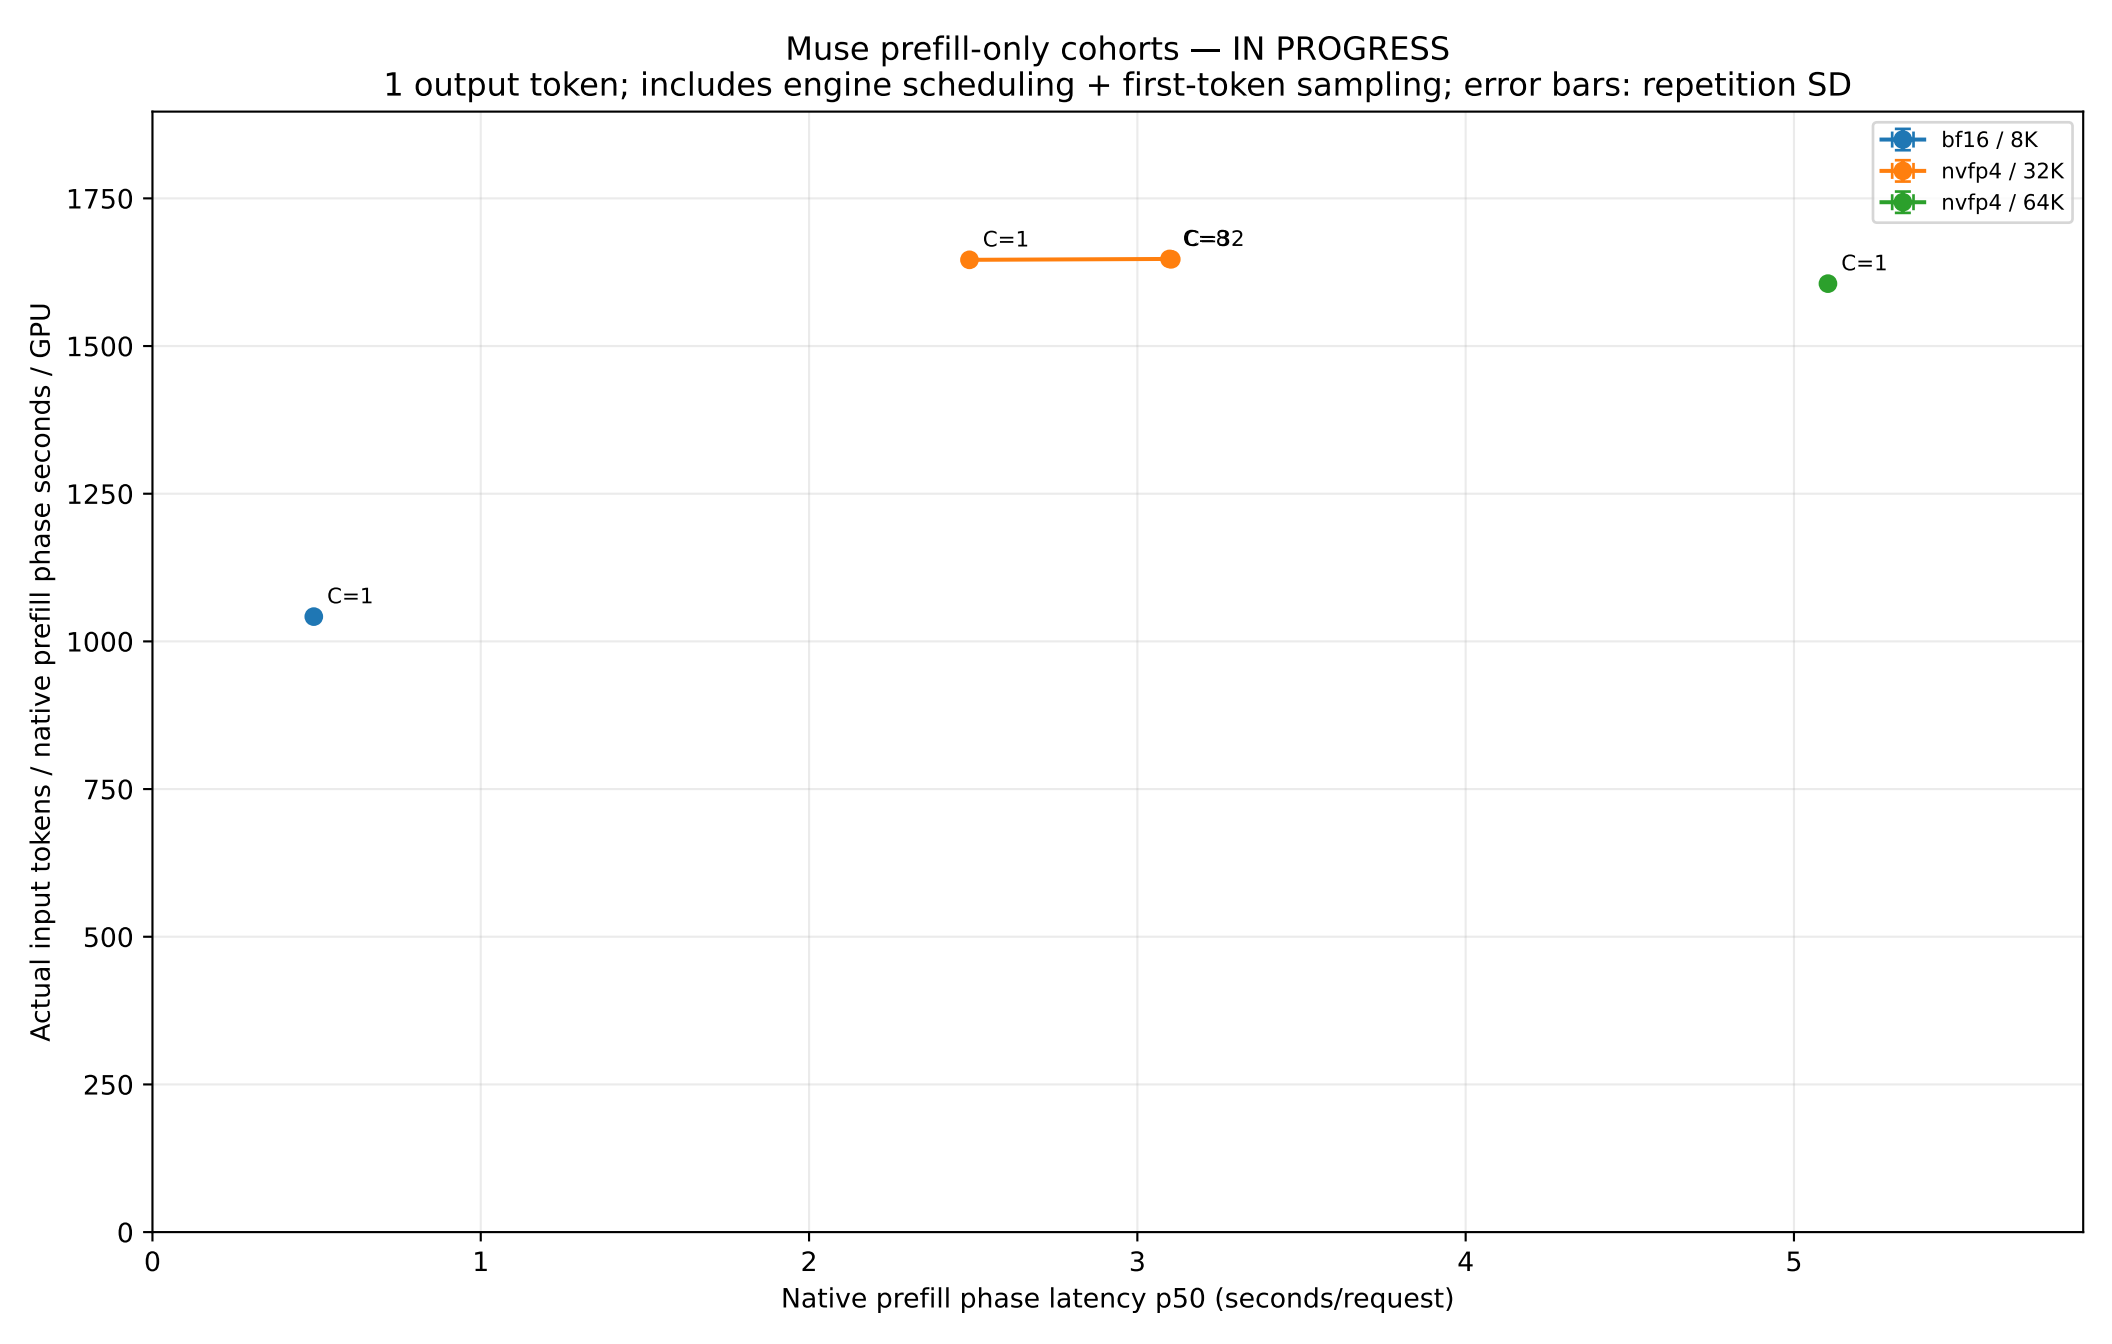
Muse prefill-only cohorts — IN PROGRESS
1 output token; includes engine scheduling + first-token sampling; error bars: repetition SD
bf16 / 8K
nvfp4 / 32K
nvfp4 / 64K
C=1
C=8
C=2
C=1
C=1
Actual input tokens / native prefill phase seconds / GPU
Native prefill phase latency p50 (seconds/request)
0
250
500
750
1000
1250
1500
1750
0
1
2
3
4
5