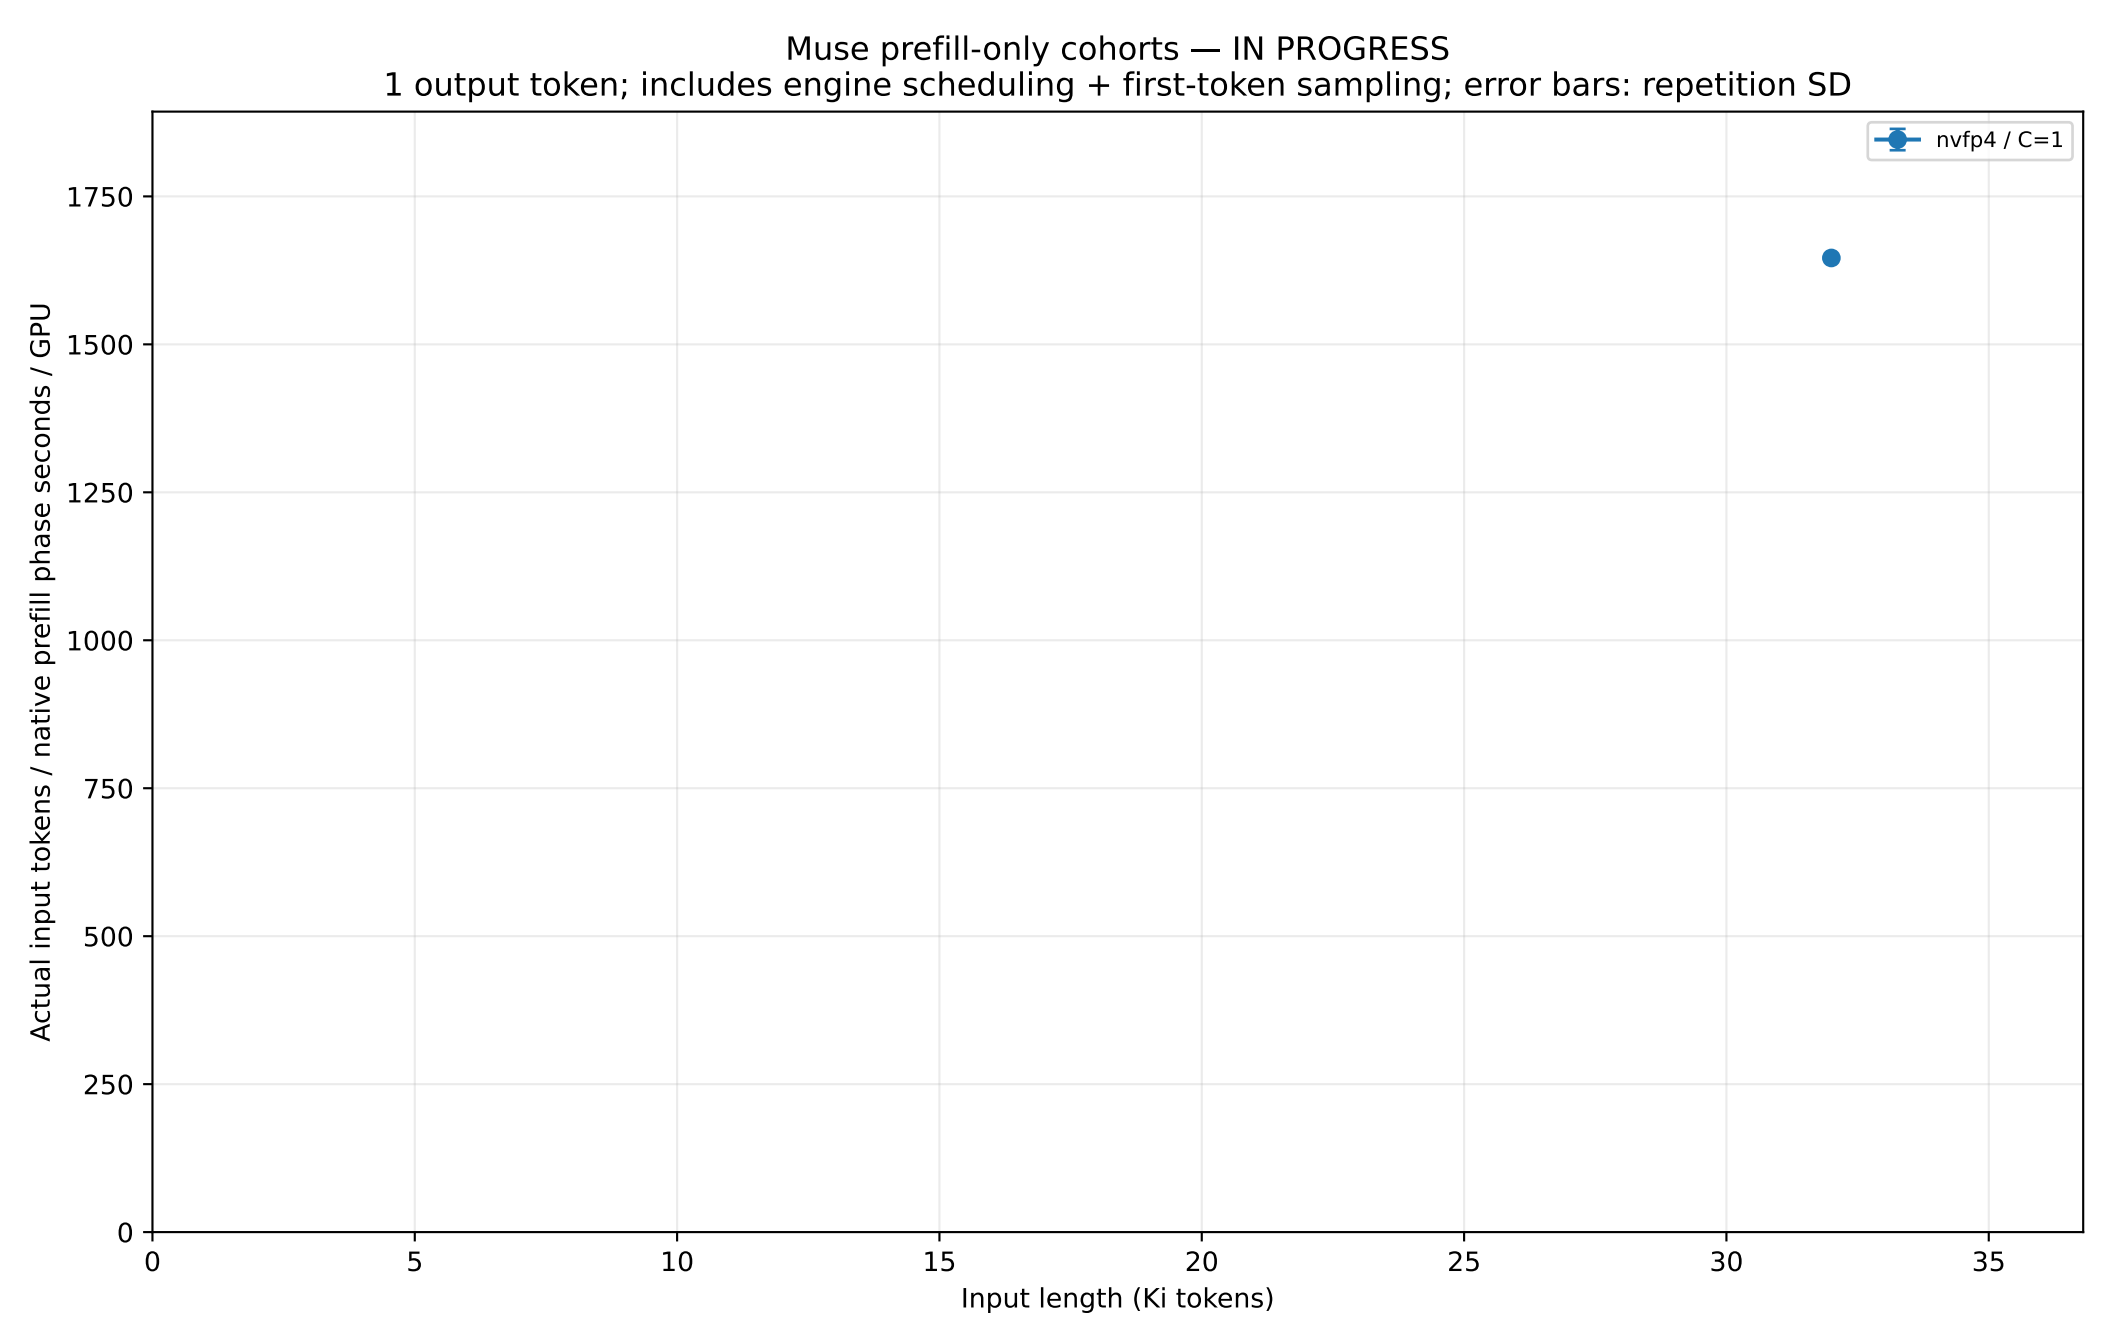
Muse prefill-only cohorts — IN PROGRESS
1 output token; includes engine scheduling + first-token sampling; error bars: repetition SD
nvfp4 / C=1
Actual input tokens / native prefill phase seconds / GPU
0
250
500
750
1000
1250
1500
1750
0
5
10
15
20
25
30
35
Input length (Ki tokens)