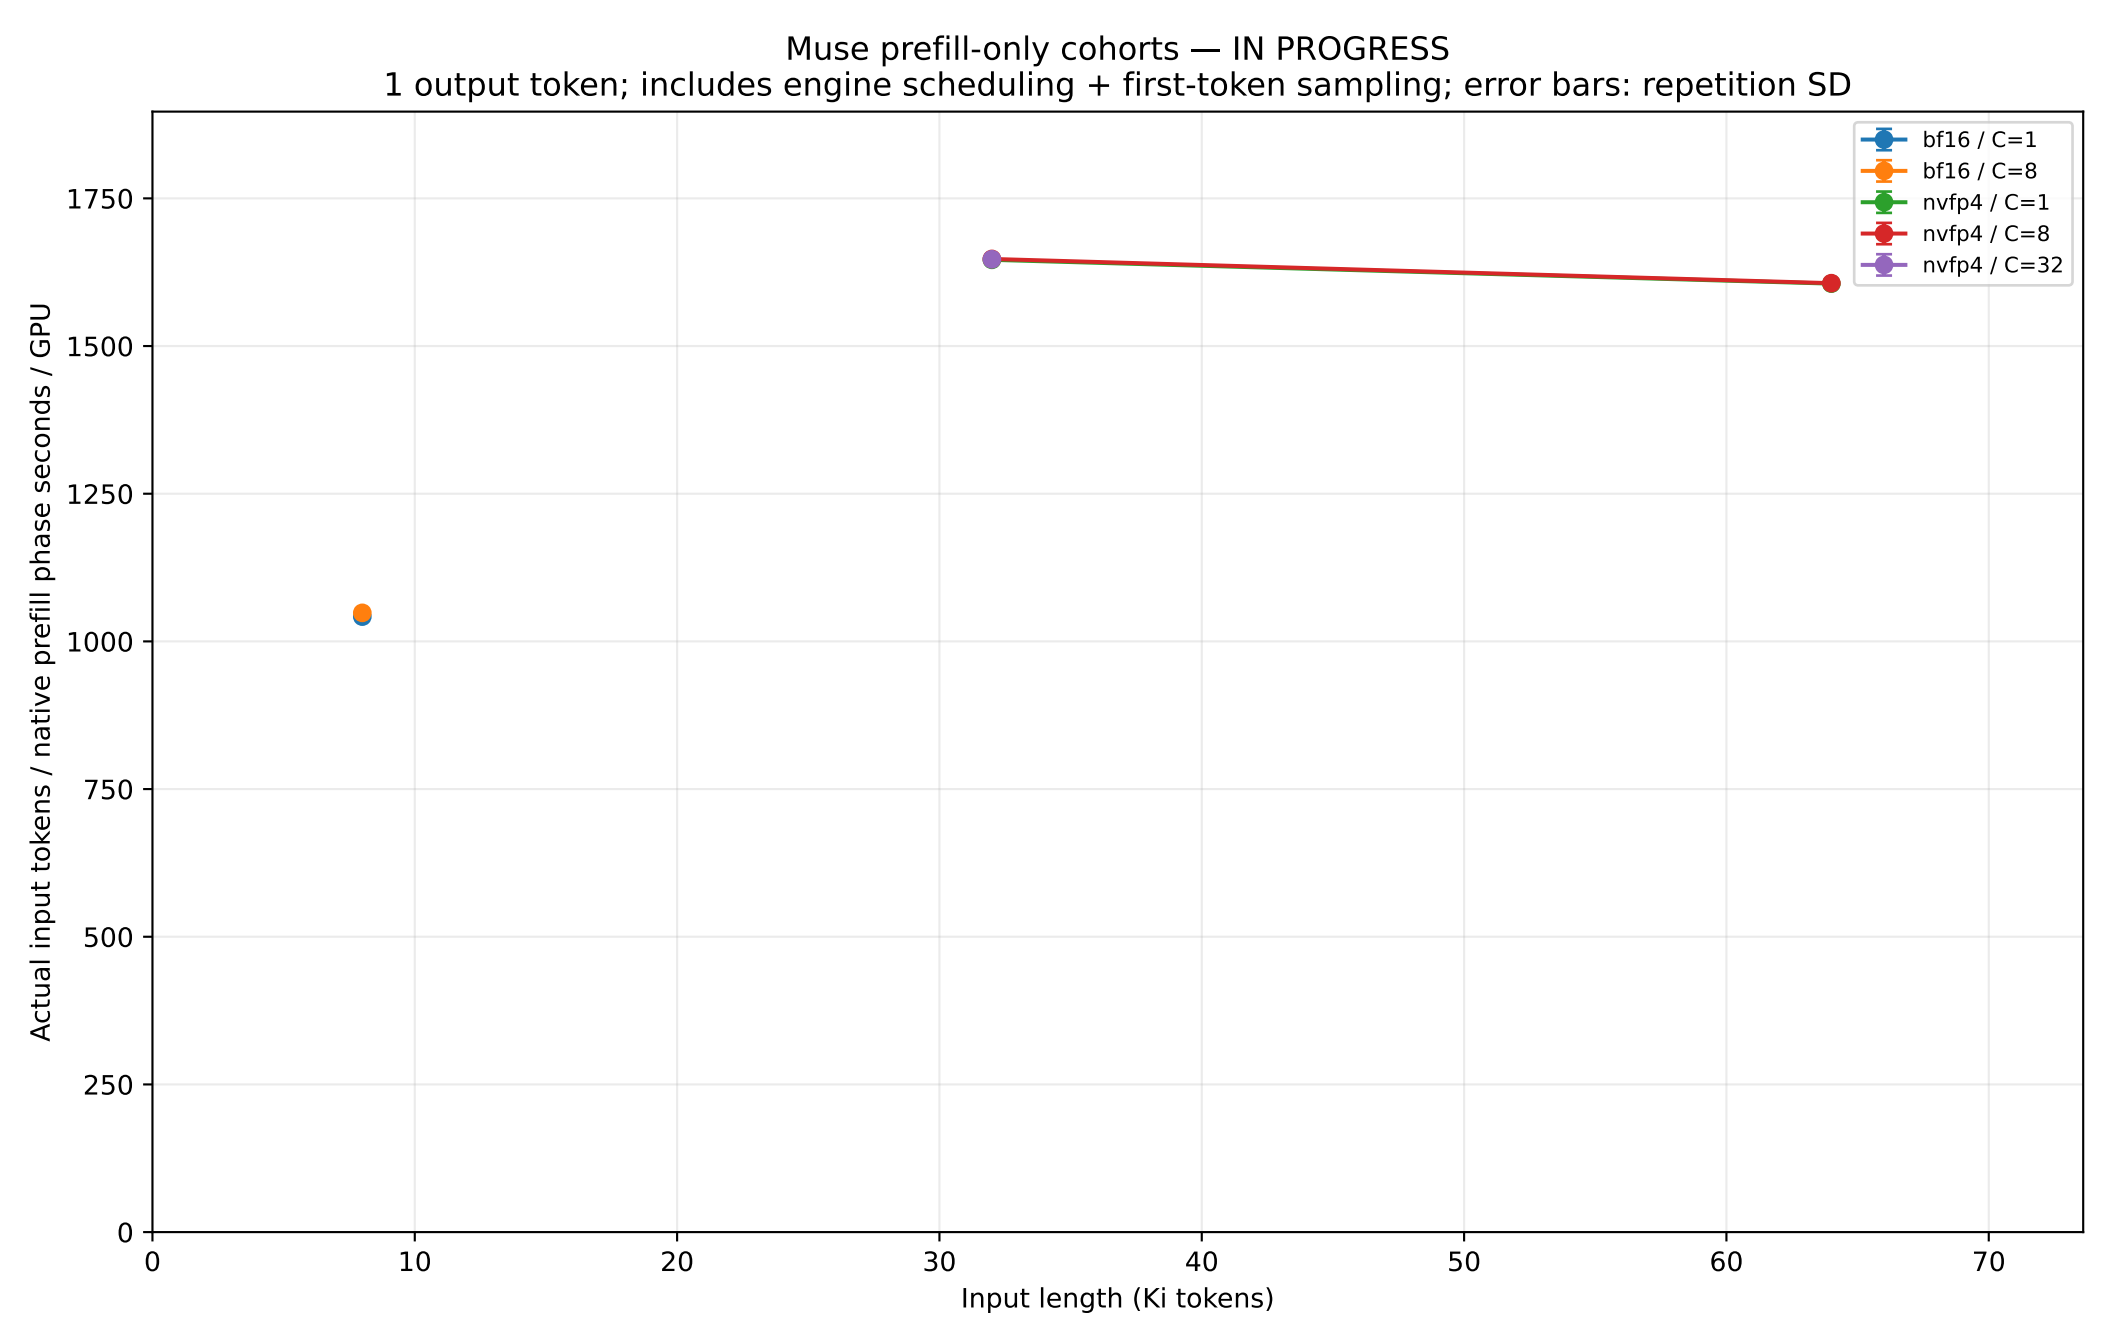
Muse prefill-only cohorts — IN PROGRESS
1 output token; includes engine scheduling + first-token sampling; error bars: repetition SD
Actual input tokens / native prefill phase seconds / GPU
Input length (Ki tokens)
0
250
500
750
1000
1250
1500
1750
0
10
20
30
40
50
60
70
bf16 / C=1
bf16 / C=8
nvfp4 / C=1
nvfp4 / C=8
nvfp4 / C=32