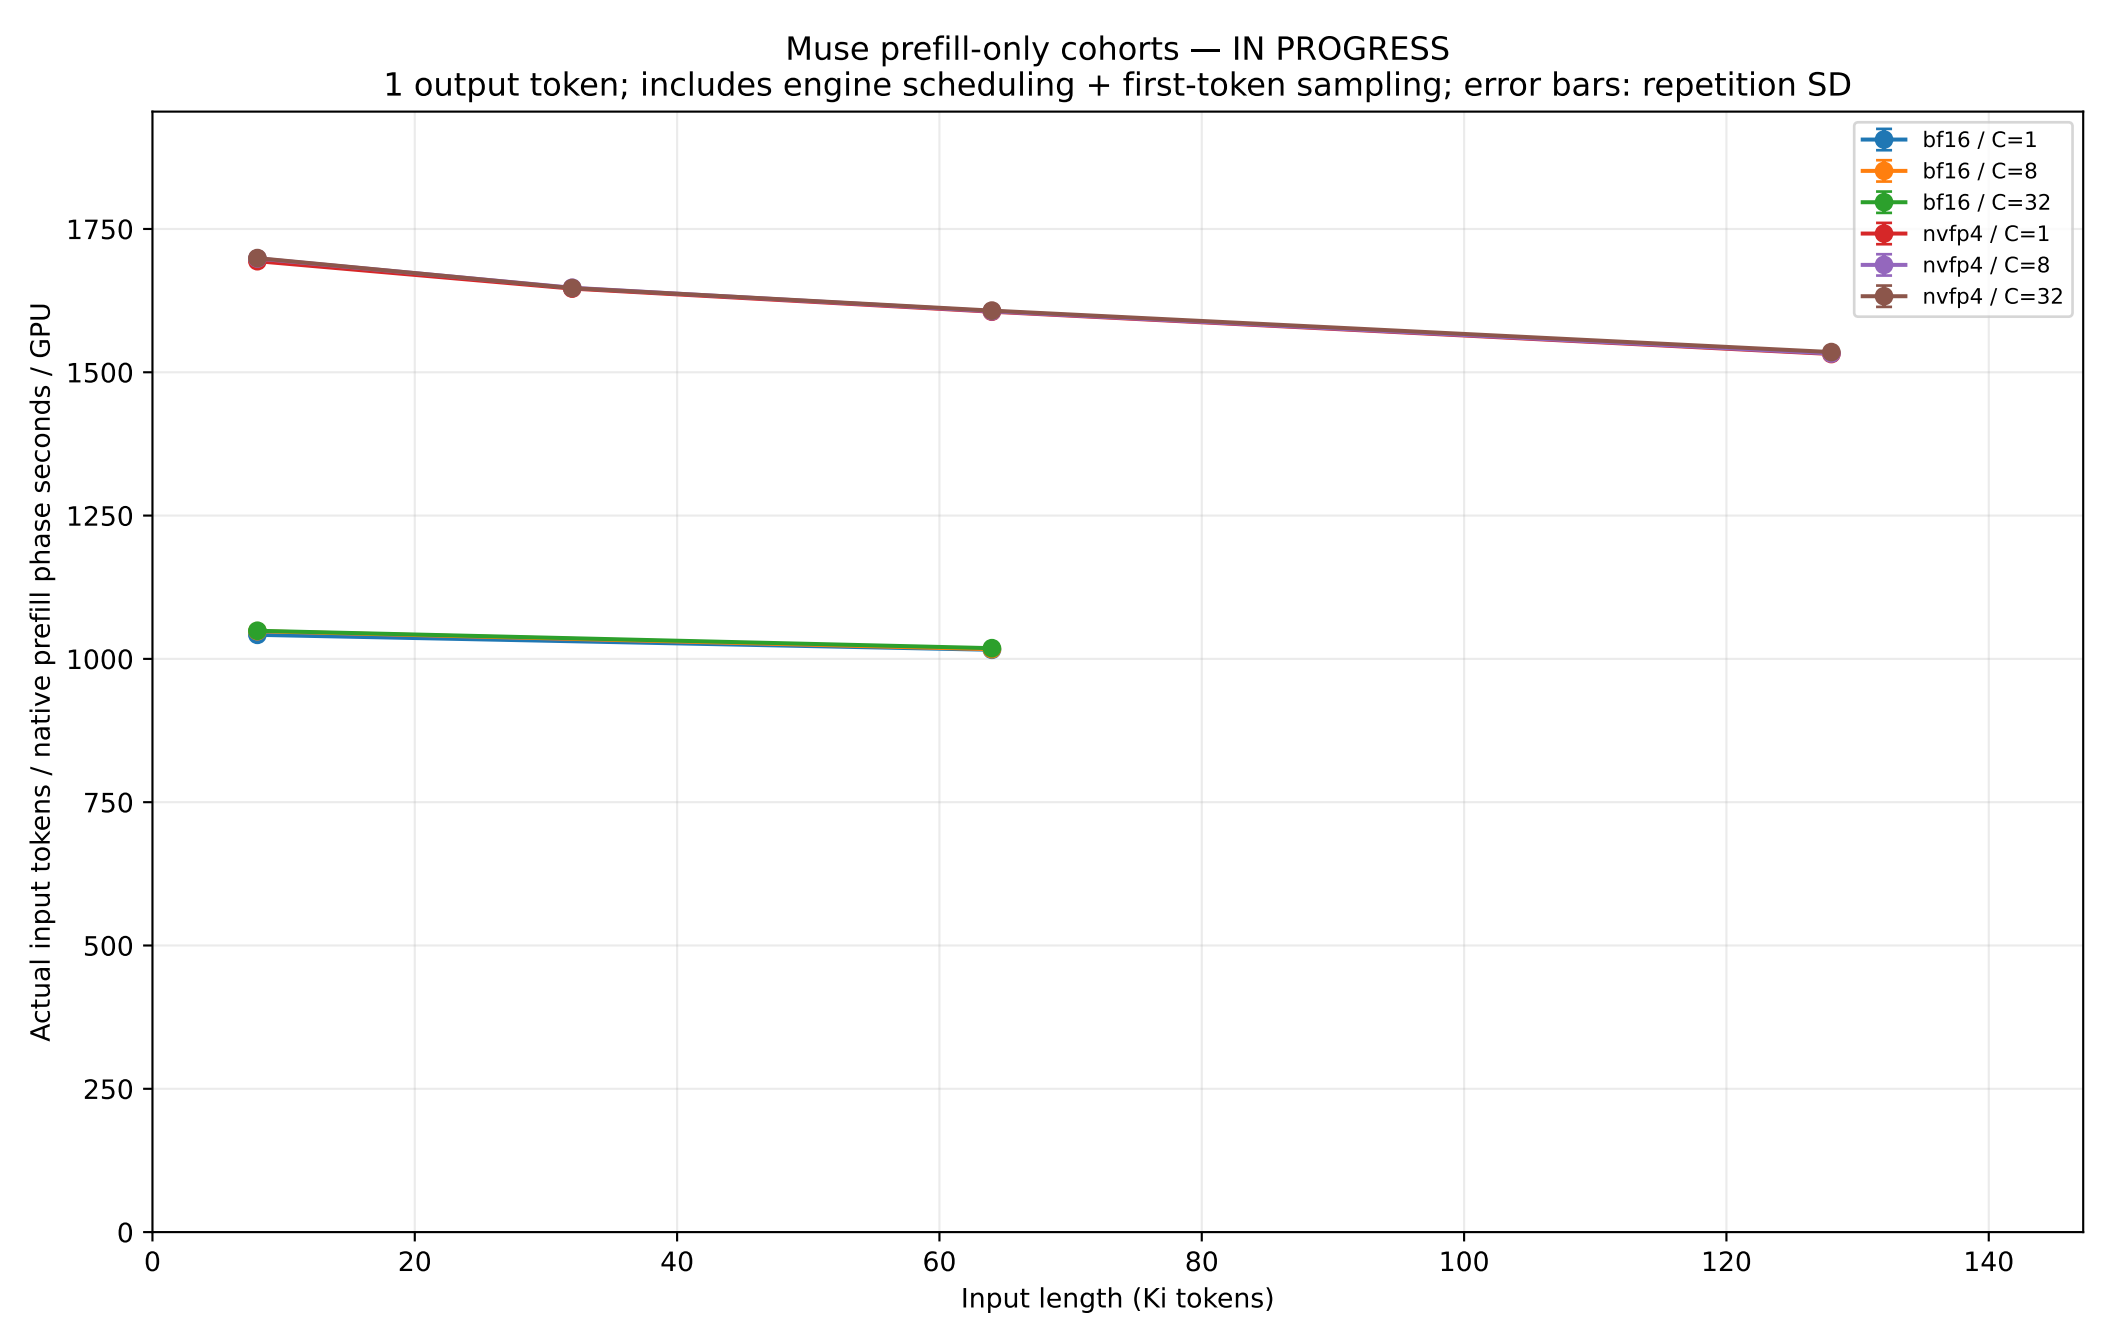
Muse prefill-only cohorts — IN PROGRESS
1 output token; includes engine scheduling + first-token sampling; error bars: repetition SD
Actual input tokens / native prefill phase seconds / GPU
Input length (Ki tokens)
bf16 / C=1
bf16 / C=8
bf16 / C=32
nvfp4 / C=1
nvfp4 / C=8
nvfp4 / C=32
0
250
500
750
1000
1250
1500
1750
0
20
40
60
80
100
120
140